



ARC309

# From Monolithic to Microservices

## Evolving Architecture Patterns in the Cloud

Adrian Trenaman  
SVP Engineering, gilt.com  
[@adrian\\_trenaman](https://twitter.com/adrian_trenaman)

Derek Chiles  
Sr. Mgr, Solutions Architecture, AWS  
[@derekchiles](https://twitter.com/derekchiles)

October 2015

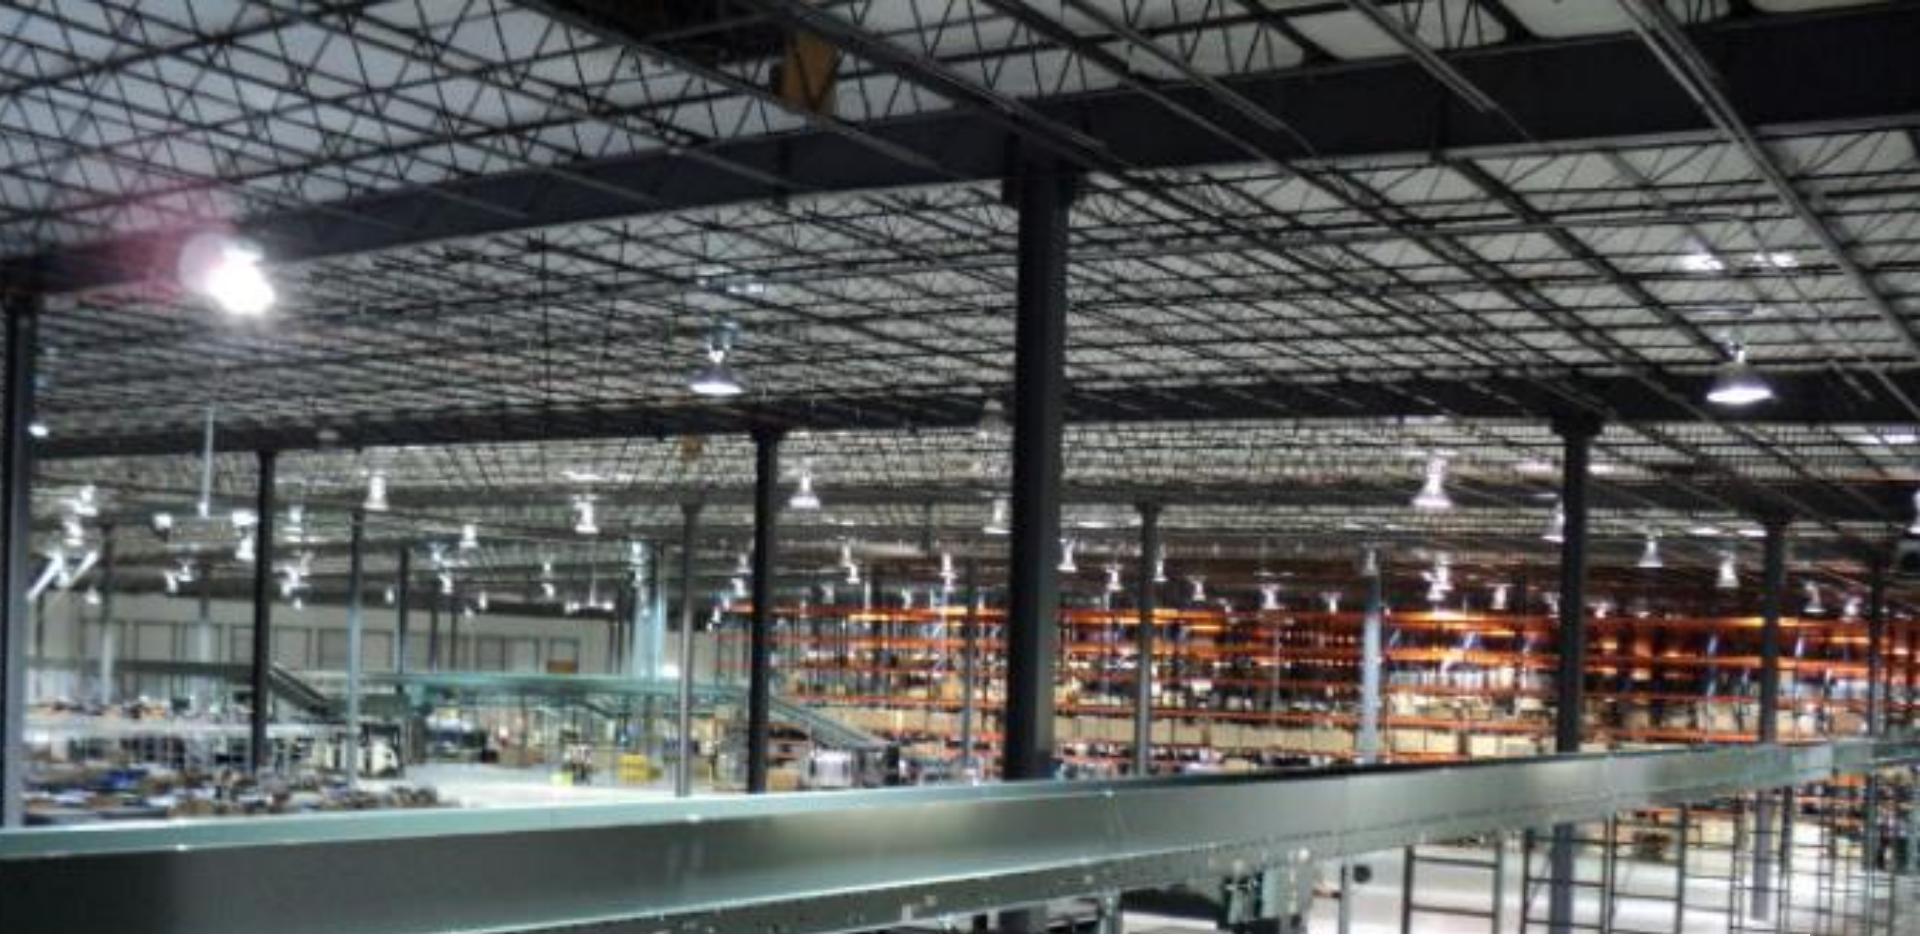
# Gilt's Journey



Gilt: Luxury designer brands at members-only prices



... we shoot the product in our studios



... we receive, store, pack, and ship ...



Shipping available to Ireland Up to 60% off the best brands.

# Wolford

Luxe hosiery, bodysuits, and soft knit layers from the legendary lingerie brand

[Shop this Sale](#)



Salvatore Ferragamo Shoes & Handbags



MYMU & More



Global Jetsetter



L'Space Swimwear & More



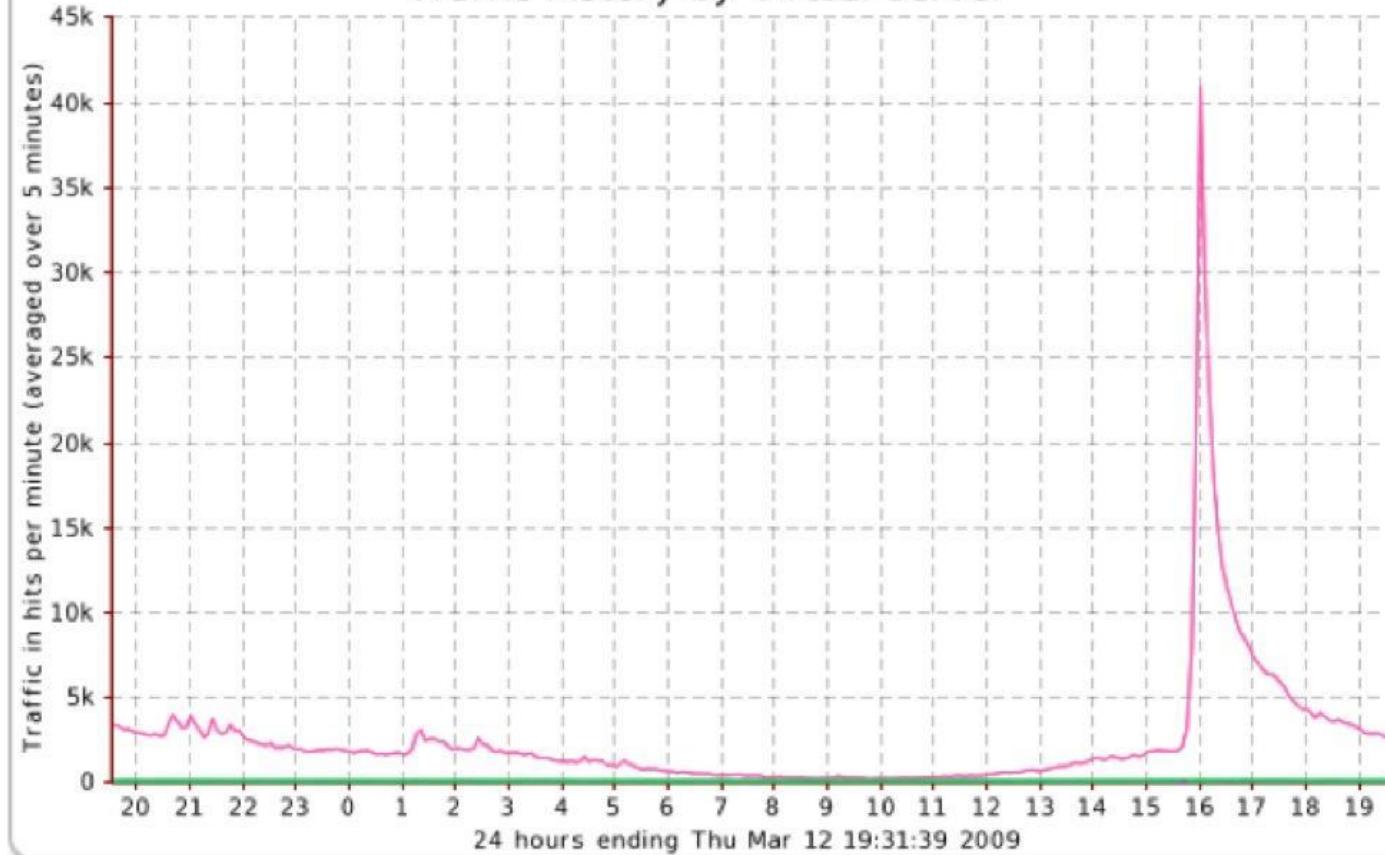
Extra 25% Off  
The World of Safavieh

... we sell every day at noon EST



... stampede!

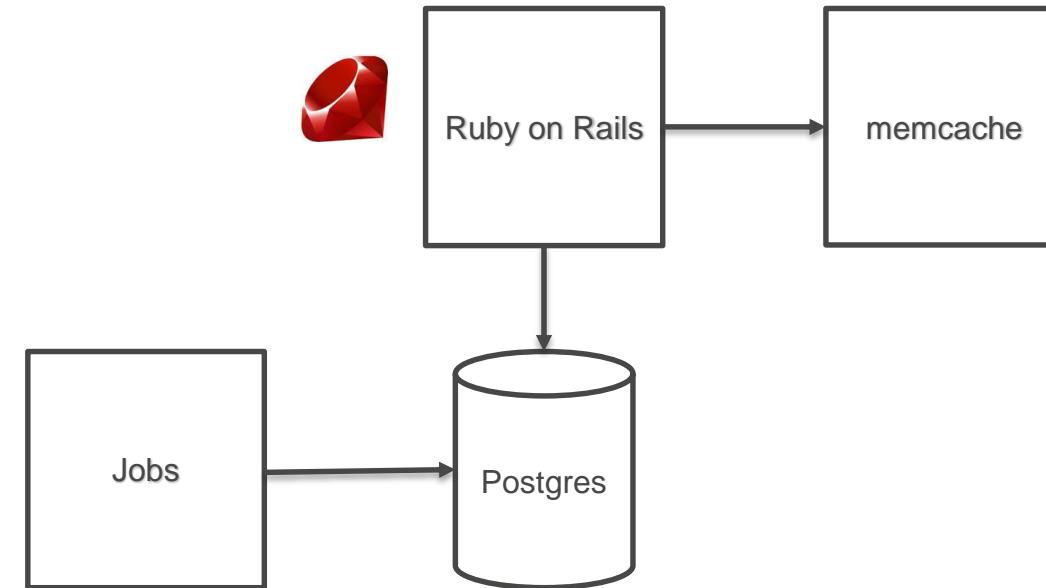
## Traffic History by Virtual Server



... this is what noon *really* looks like.

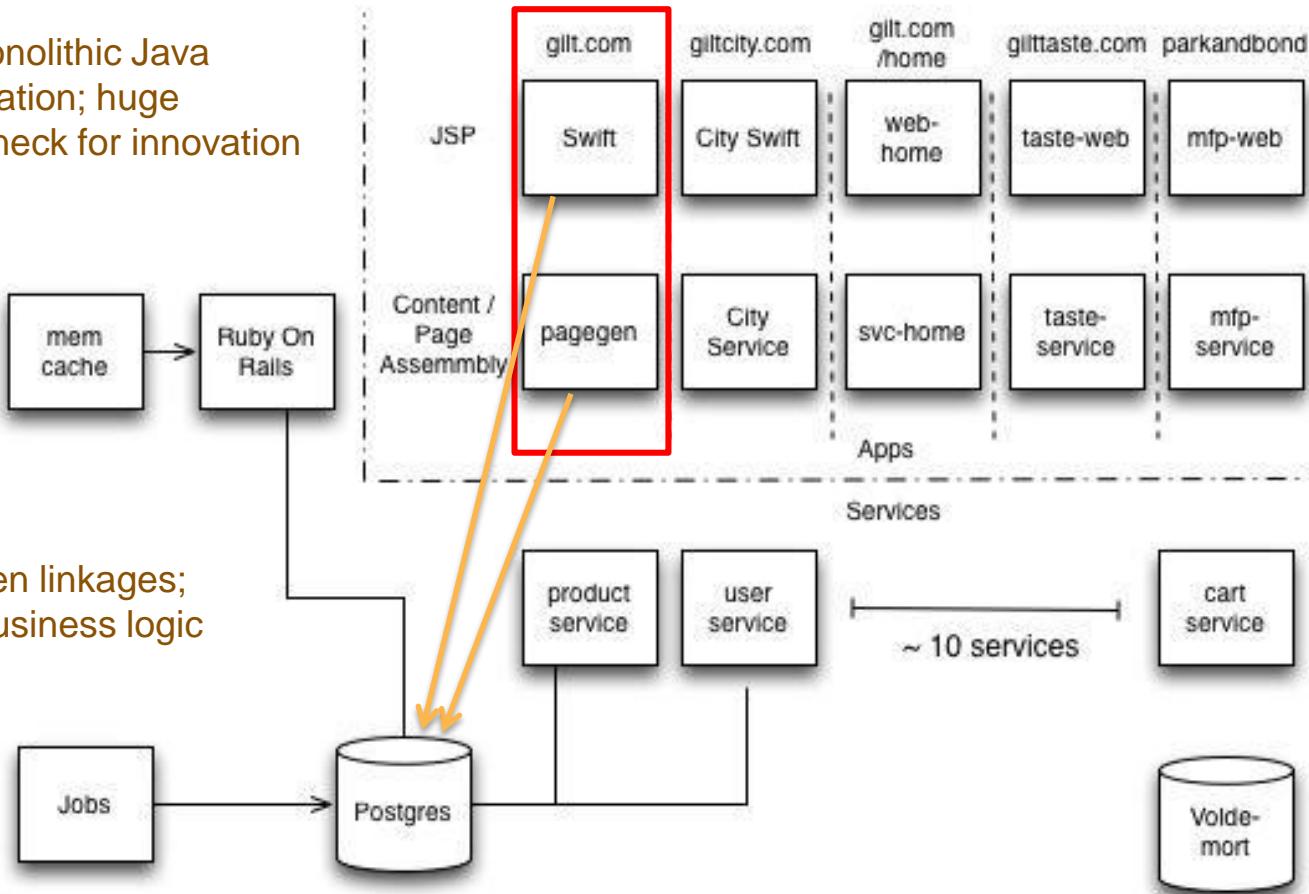


# **How It All Started...**



From Rails to Riches — 2007: A Ruby on Rails monolith

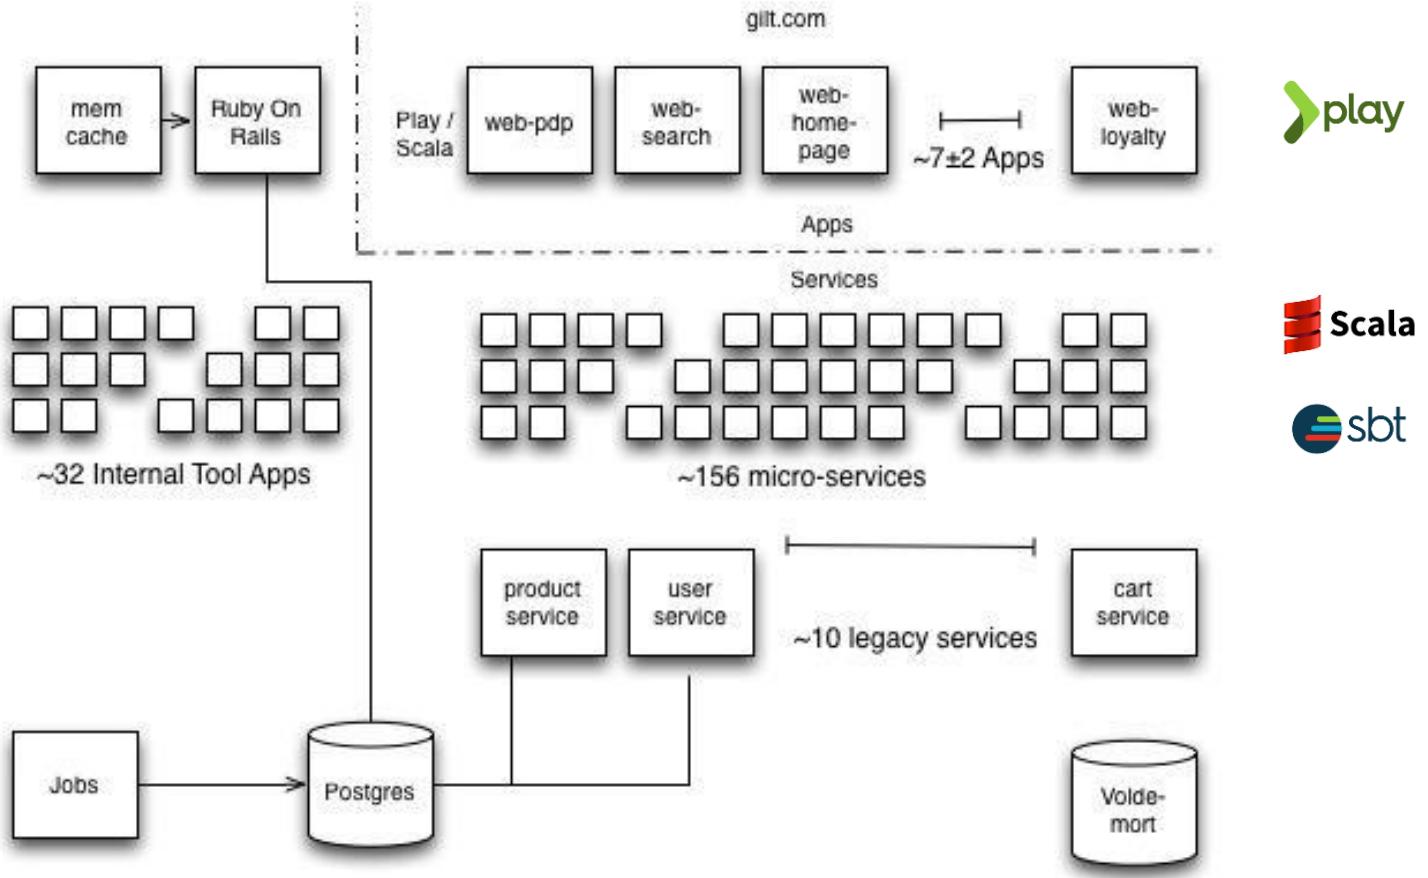
(3) Monolithic Java application; huge bottleneck for innovation



2011: Java, loosely-typed, monolithic services

“How can we arrange our teams around strategic initiatives? How can we make it fast and easy to get to change to production?”

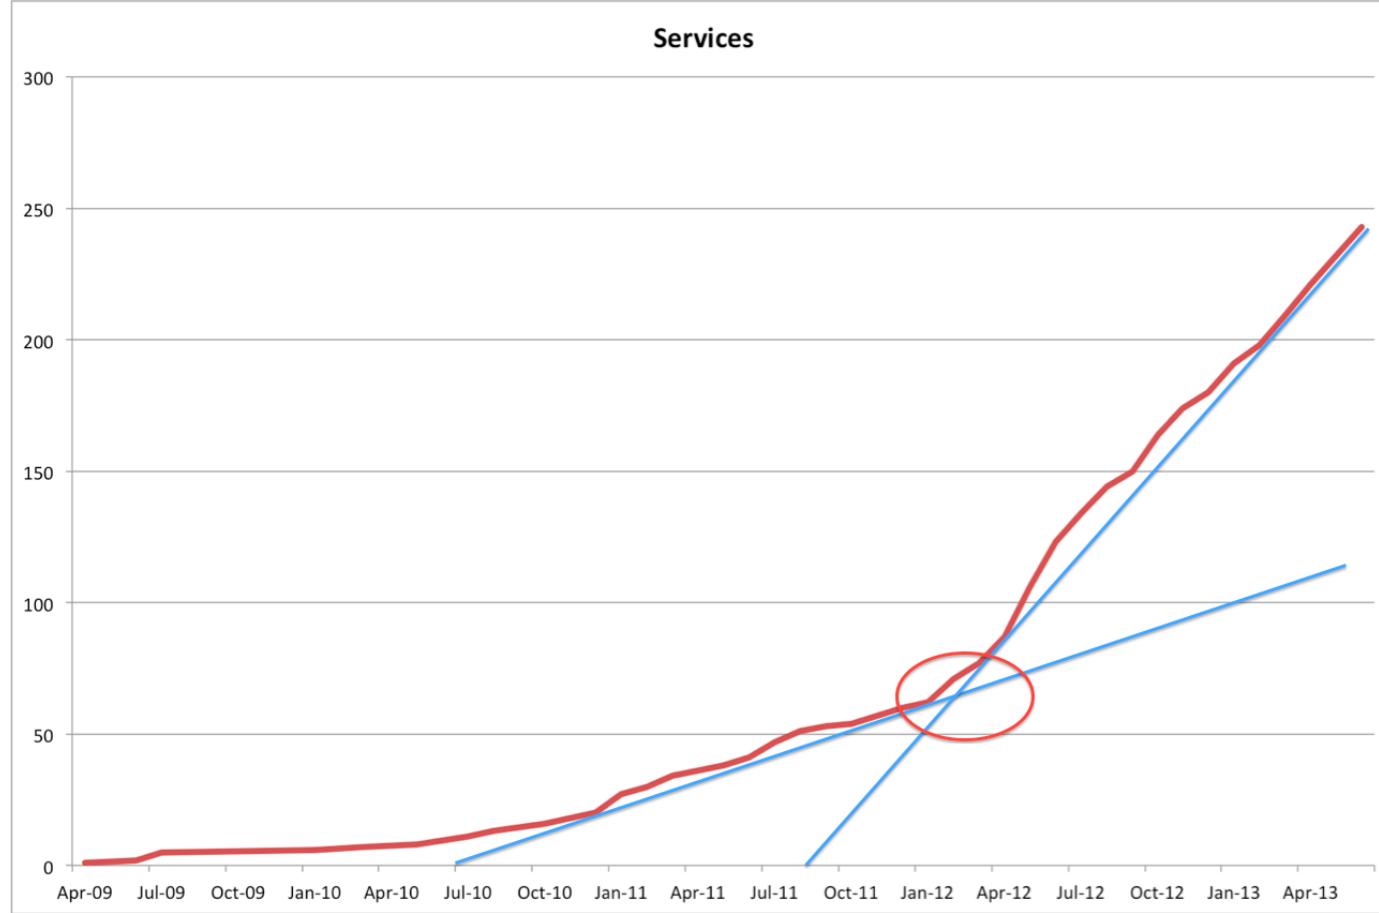
Enter: μ-services



2015: LOSA (Lots of Small Apps) & Microservices

# Driving Forces Behind Gilt's *Emergent* Architecture

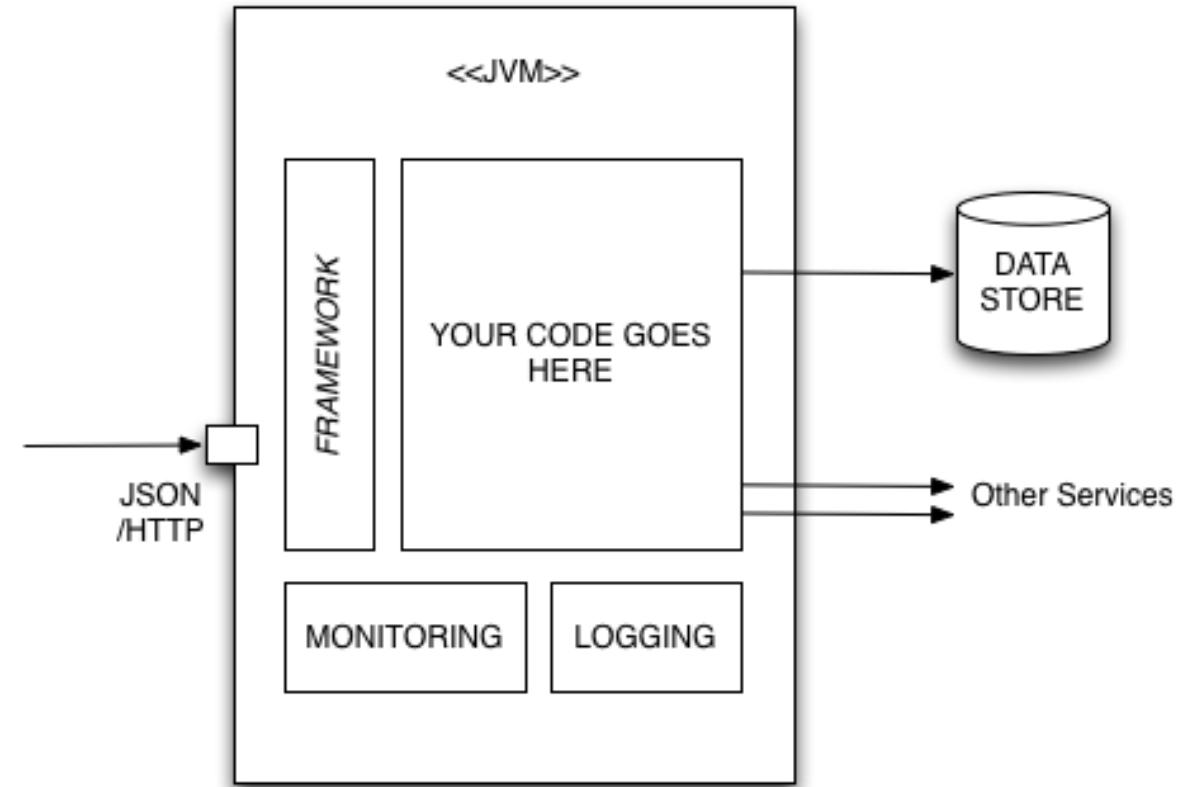
- Team autonomy
- Voluntary adoption (tools, techniques, processes)
- KPI / goal-driven initiatives
- Failing fast and openly
- Being open and honest, even when it's difficult



Service growth over time: point of inflection === Scala.

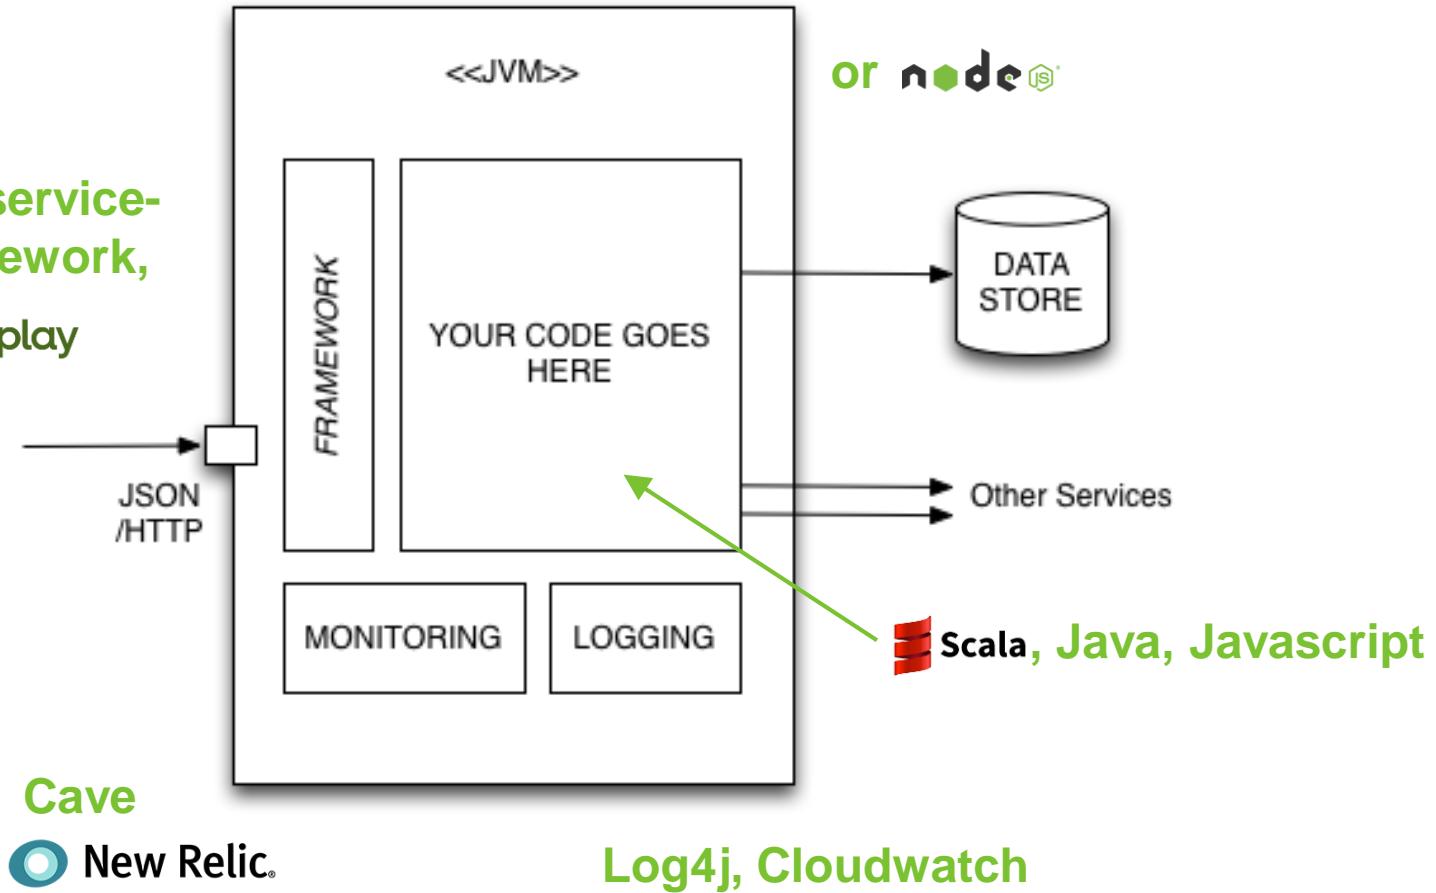
# service

# What are all these services doing?



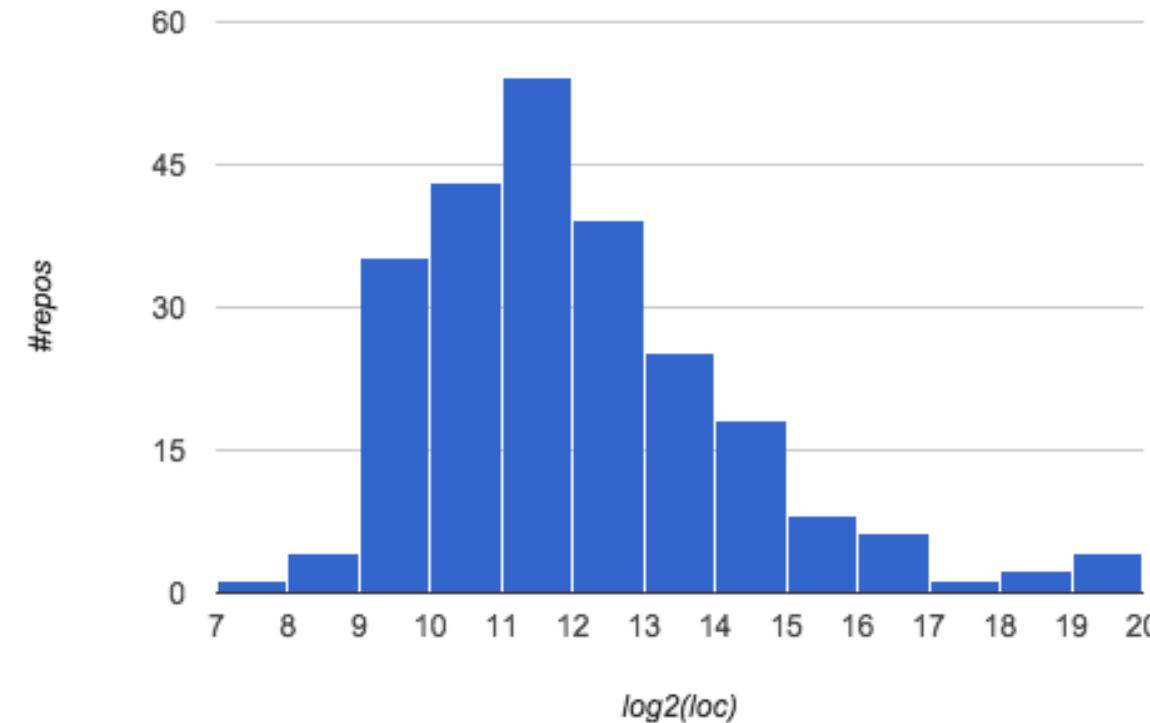
## Anatomy of a service

gilt-service-framework,



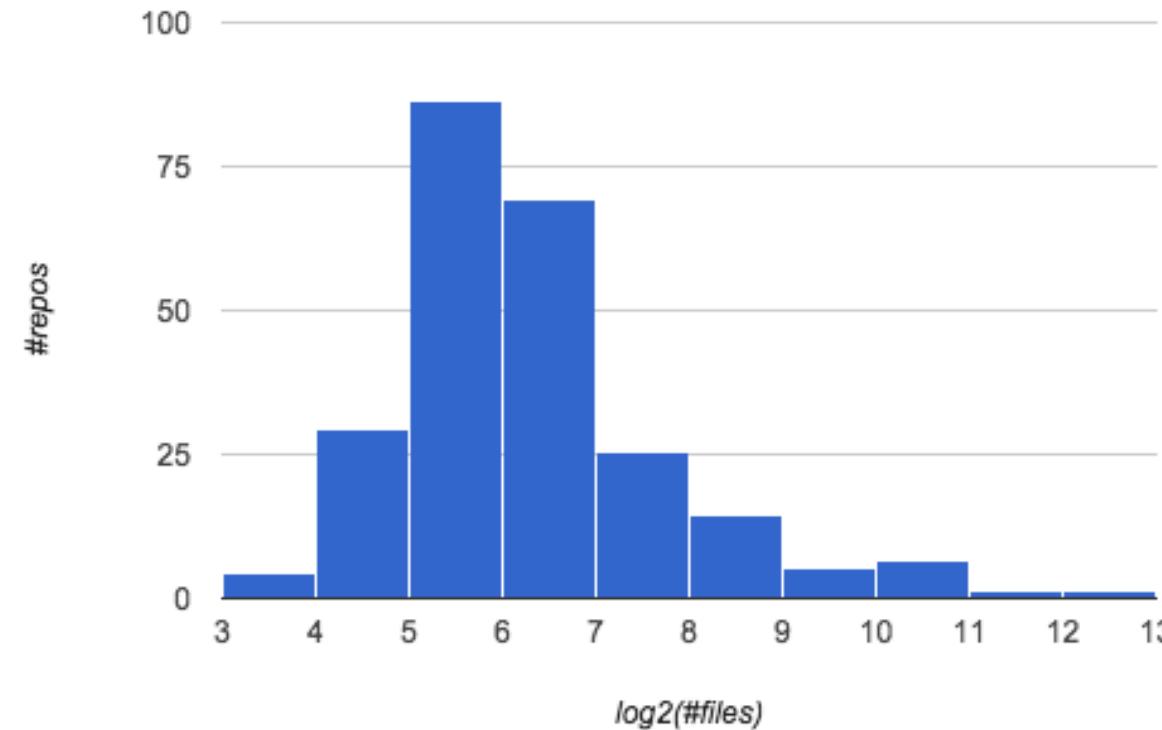
Anatomy of a gilt service – typical choices

## LOC per service (logarithmic)

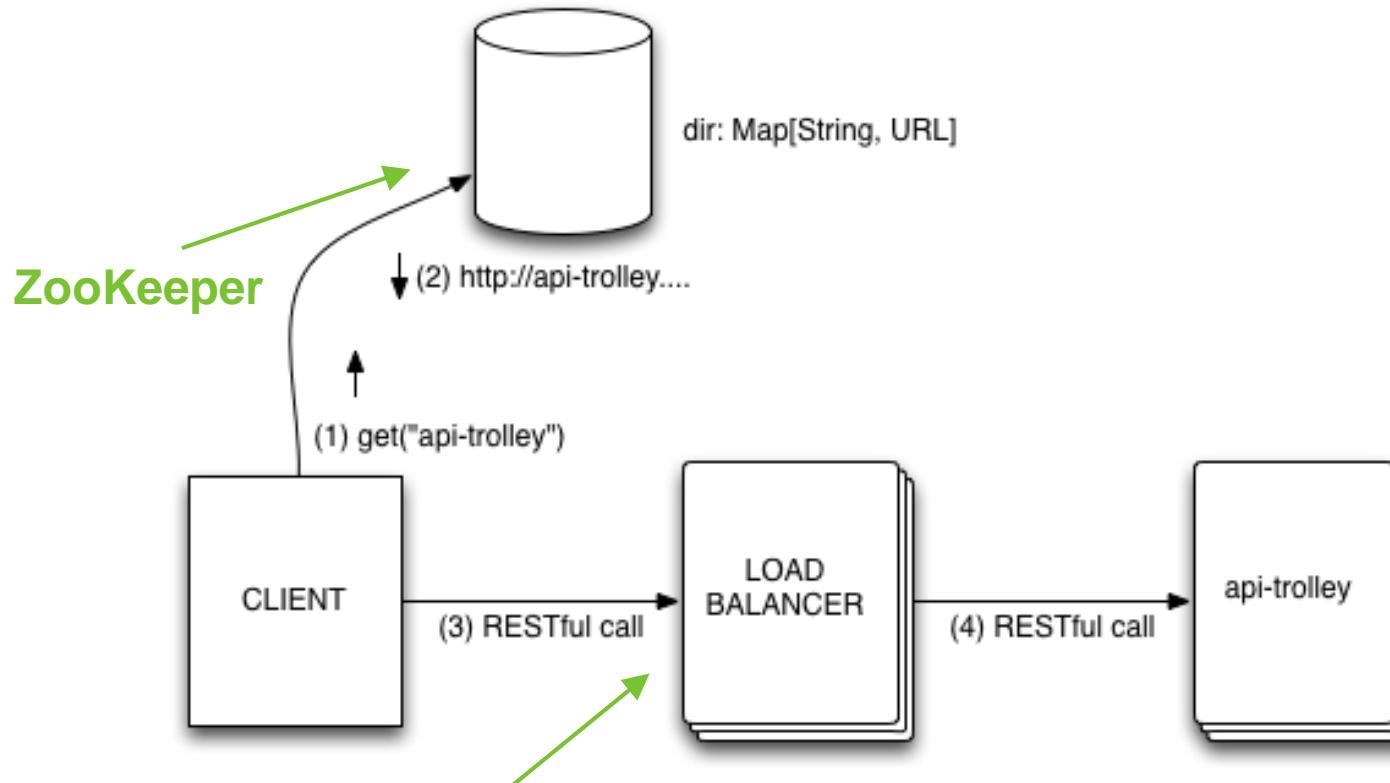


Lines of code per service (logarithmic scale)

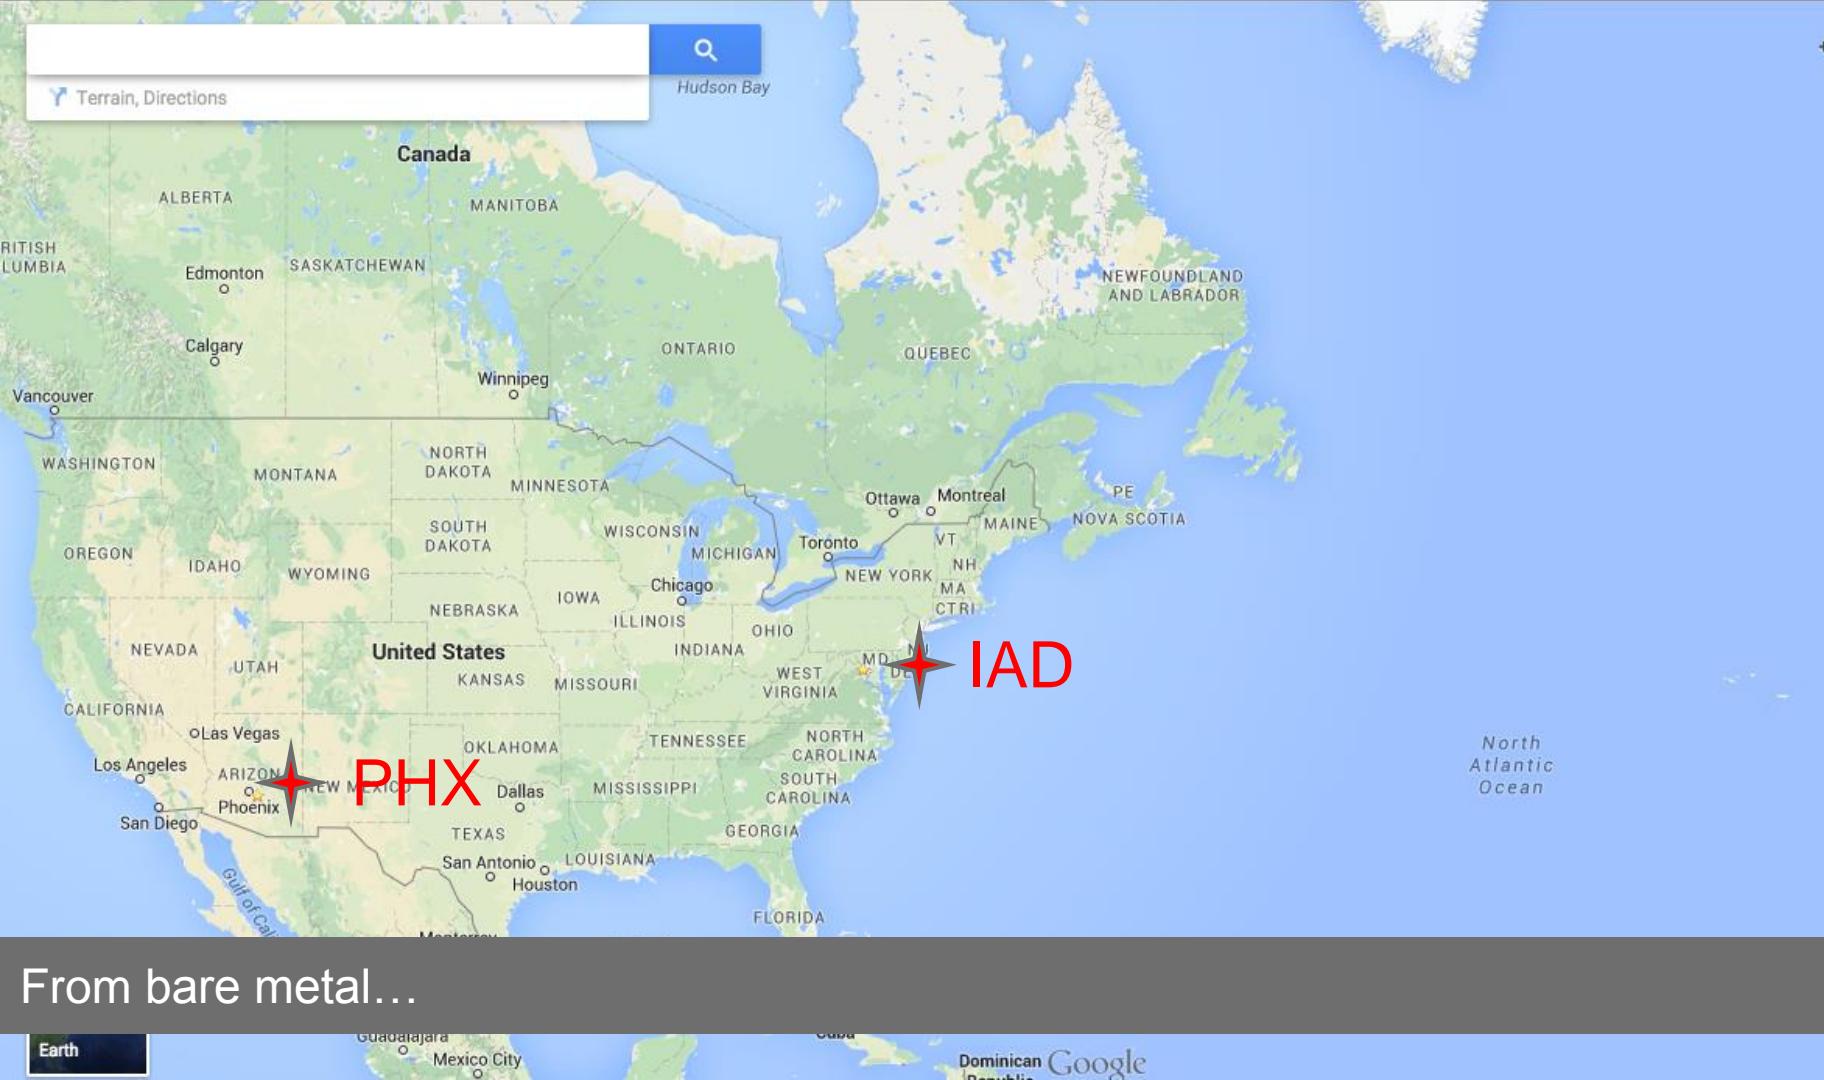
## Source files per service (logarithmic)



# source files per service (includes build, config, xml, Java, Scala, Ruby...)



Service discovery: straightforward



From bare metal...



... to vapour.

## Existing Data Centre

(1) Deploy to VPC

Dual 10-Gb direct connect line, 2-ms latency.

“Legacy VPC”

(2) “Department” accounts for elasticity & DevOps

Common

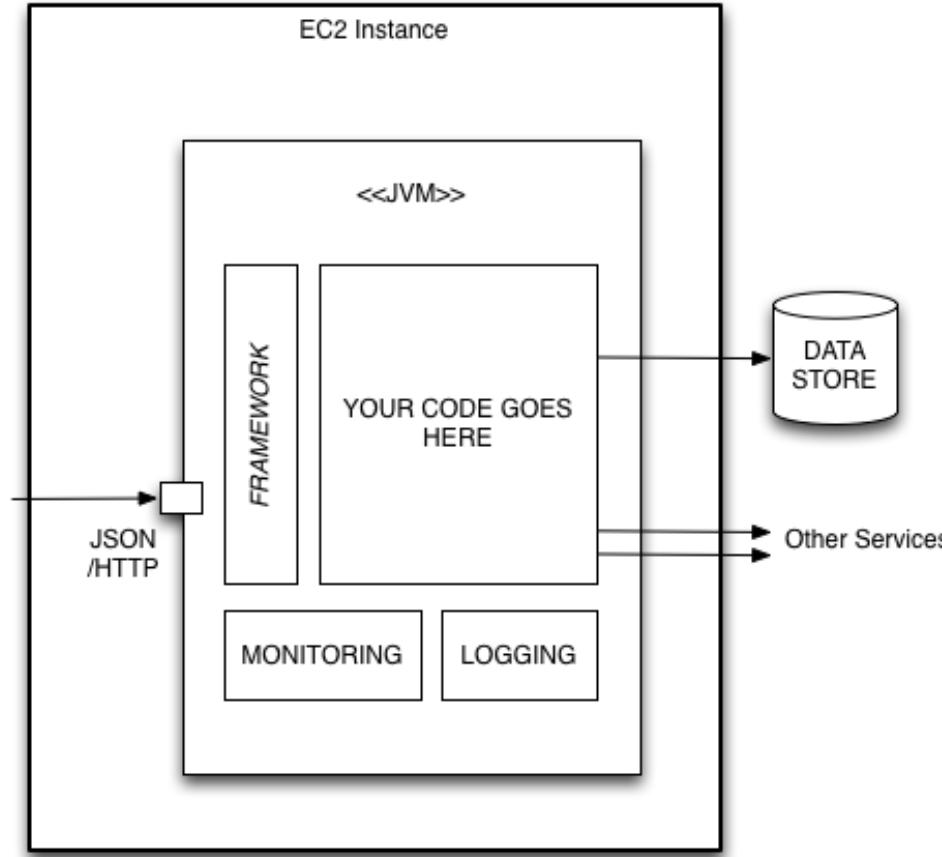
Mobile

Person-  
alisation

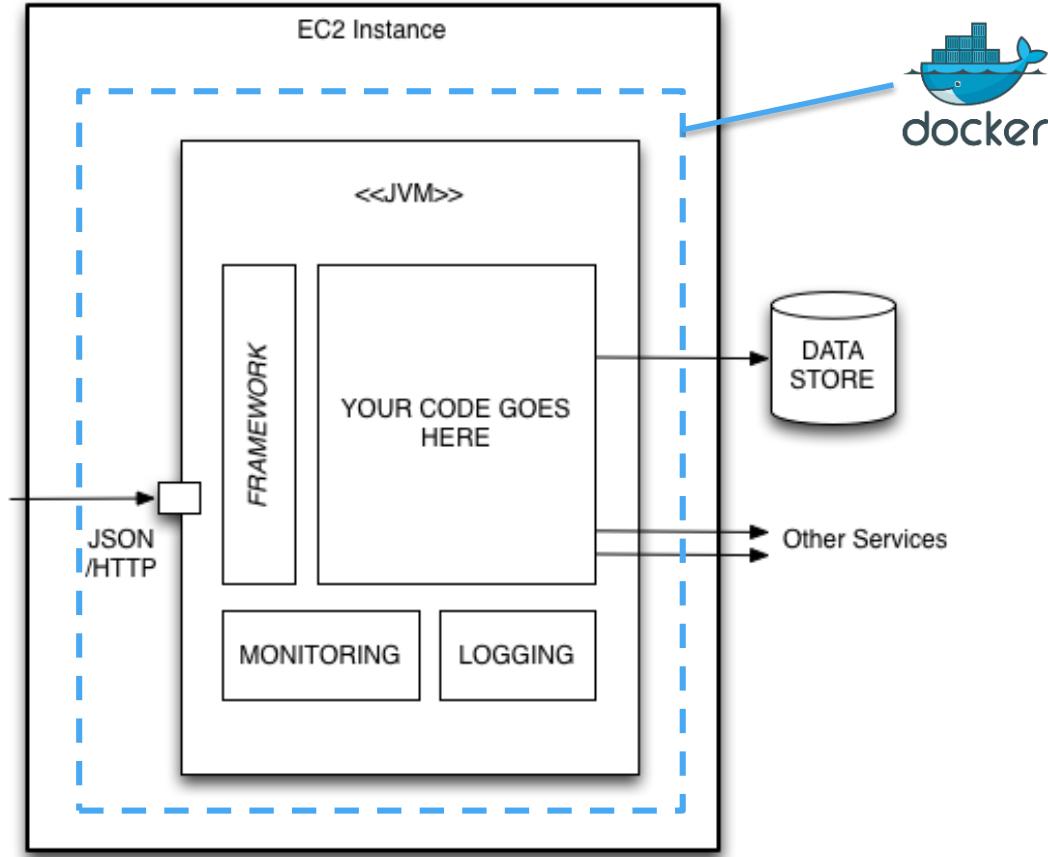
Admin

Data

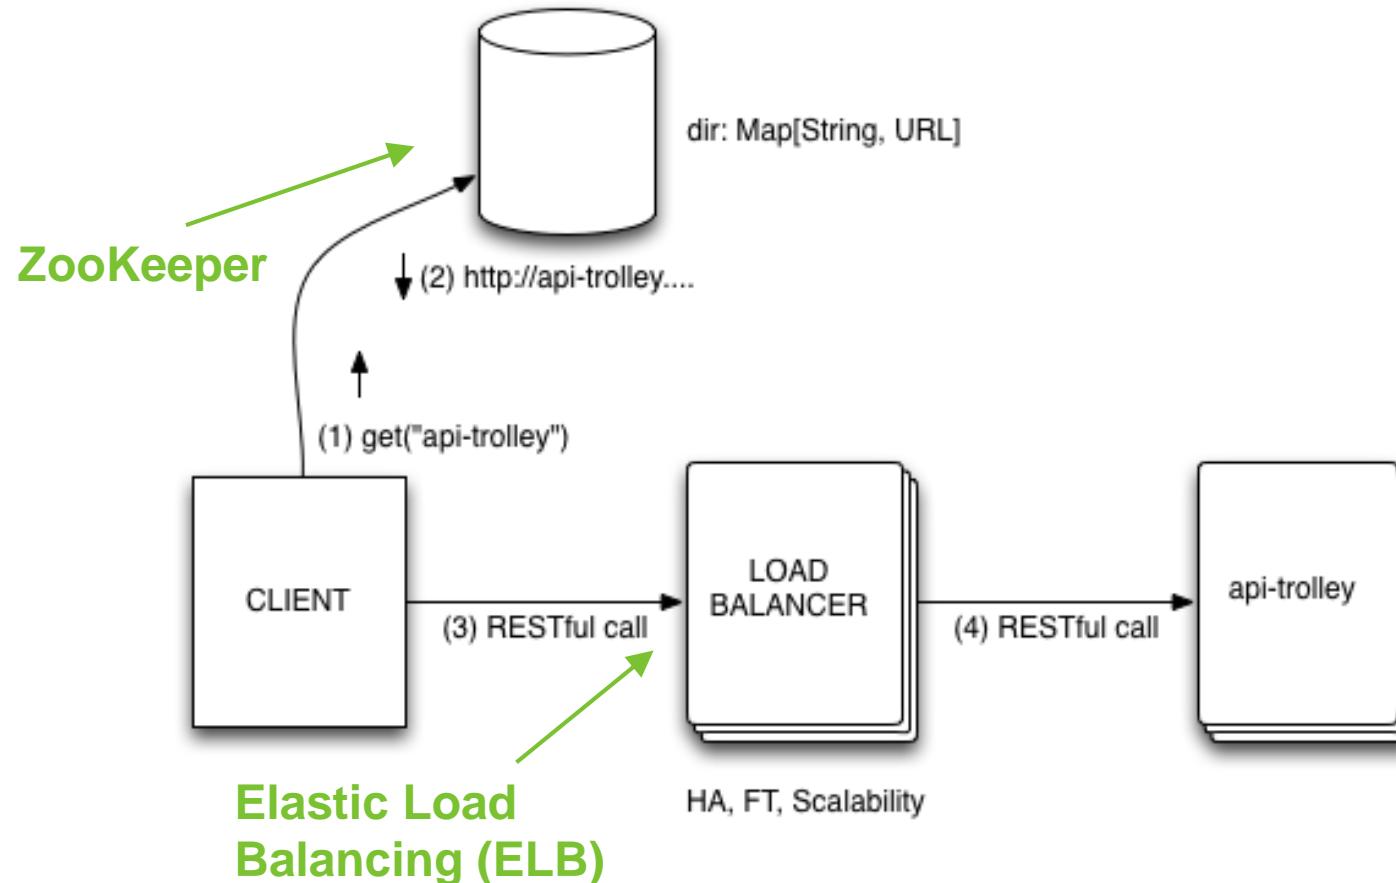
Lift-and-shift + elastic teams



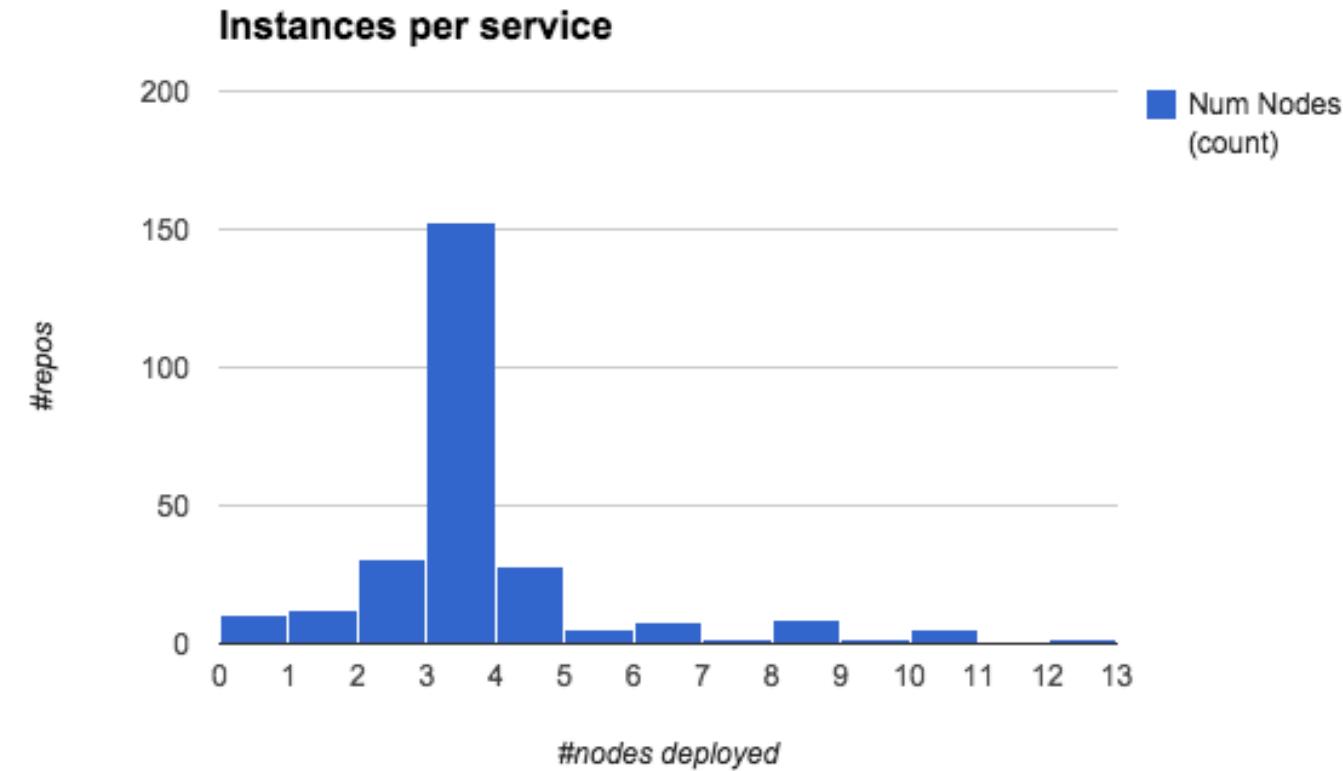
Single-tenant deployment: one service per EC2 instance



Reproducible, immutable deployments: Docker

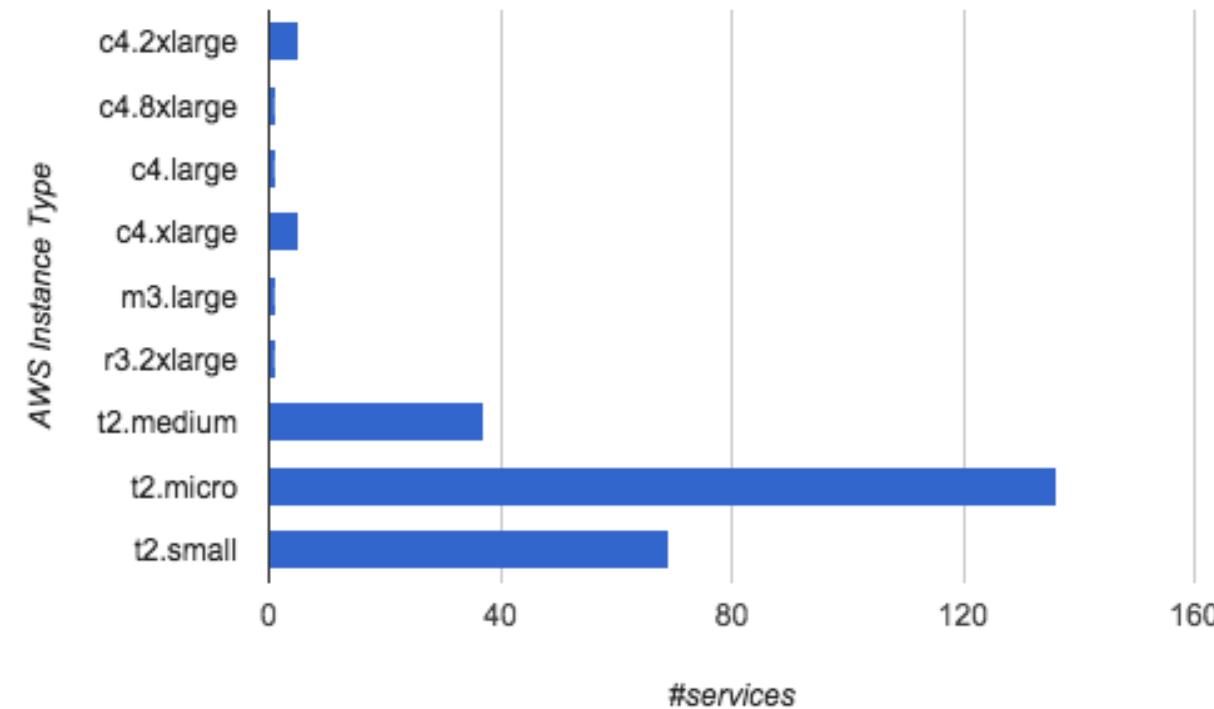


Service discovery: new services use ELB



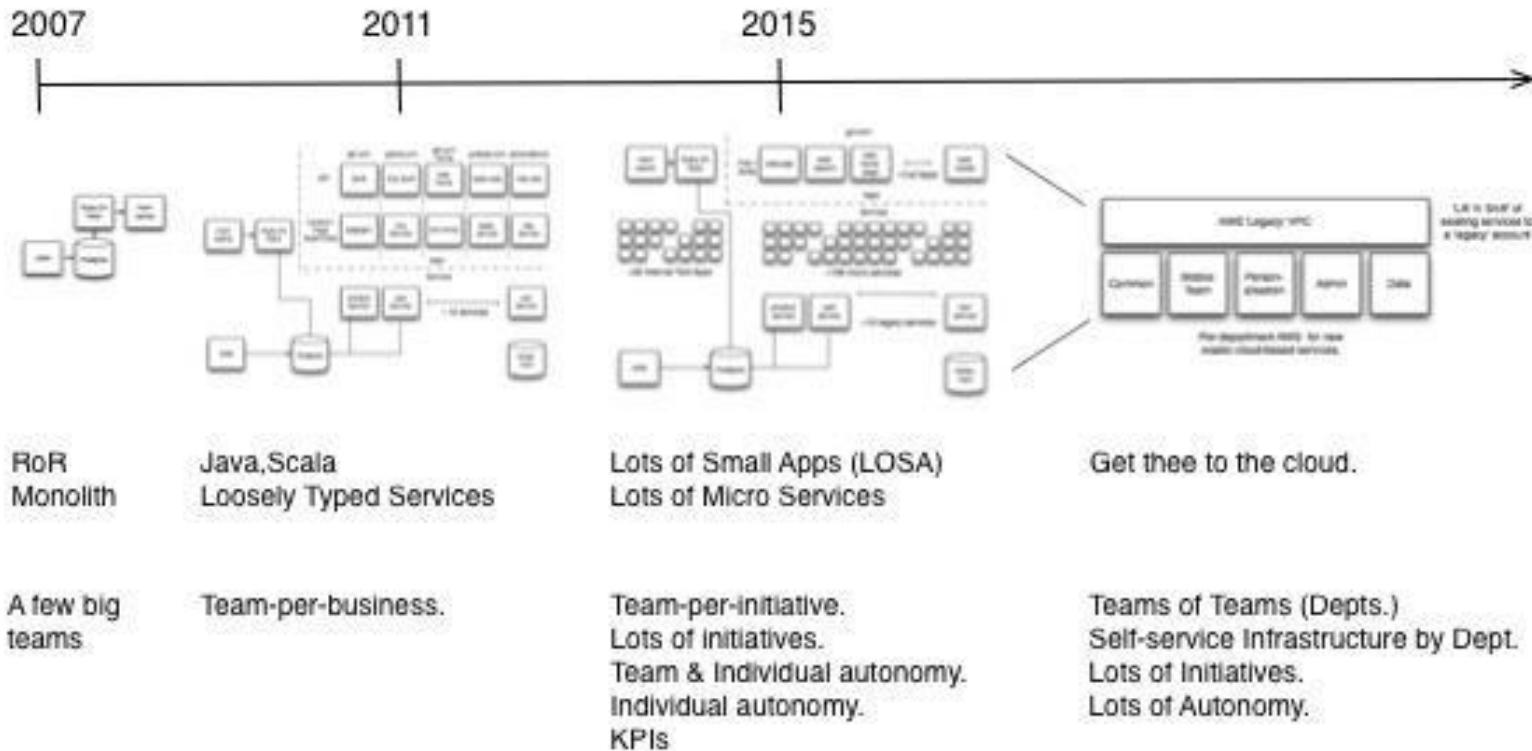
# running instances per service: “rule of three” (previously “rule of four”)

## AWS Instance Size



EC2 instance sizing: lots of small instances

## GILT: Evolution of a Micro-Services Architecture in a \$1B startup



Evolution of architecture and tech organization

# We (heart) μ-services

- Lessen dependencies between teams: faster code-to-prod
- Lots of initiatives in parallel
- Your favourite <tech/language/framework> here
- Graceful degradation of service
- Disposable code: easy to innovate, easy to fail and move on.

# We (heart) cloud

- Do DevOps in a *meaningful* way.
- Low barrier of entry for new tech (Amazon DynamoDB, Amazon Kinesis,...)
- Isolation
- Cost visibility
- Security tools (IAM)
- Well documented
- Resilience is easy
- Hybrid is easy
- Performance is great



# **Common Challenges and Patterns**

## Monolithic

- Simple deployments
- Binary failure modes
- Inter-module refactoring
- Technology monoculture
- Vertical scaling

## Microservices

- Partial deployments
- Graceful degradation
- Strong module boundaries
- Technology diversity
- Horizontal scaling

# Common Challenges and Patterns

- Organization
- Discovery
- Data management
- Deployment
- I/O explosion
- Monitoring



Organization

# Monolithic Ownership

Organized on technology capabilities



UI Team

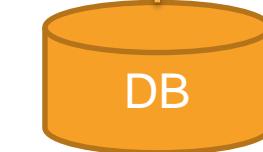


App Logic Team



DBA Team

Organizational Structure



Application Architecture

# Microservices Ownership

Organized on business responsibilities



Login  
Registration  
Order

Accounts team



Personalization

Personalization team



Mobile

Mobile team

# Microservices Ownership

- Requirements
- Technology selection
- Development
- Quality
- Deployment
- Support

# How to Be a Good Citizen (Service Consumer)

- Design for failure
- Expect to be throttled
- Retry w/ exponential backoff
- Degrade gracefully
- Cache when appropriate

# How to Be a Good Citizen (Service Provider)

- Publish your metrics
- Protect yourself
- Keep your implementation details private
- Maintain backwards compatibility

# Amazon API Gateway

- Throttling (global and per-method)
- Caching (with TTLs and invalidation)
- Monitoring (RPS, latency, error rate)
- Versioning
- Authentication



Discovery

# Use DNS

## Convention-based naming

<service-name>-<environment>.domain.com

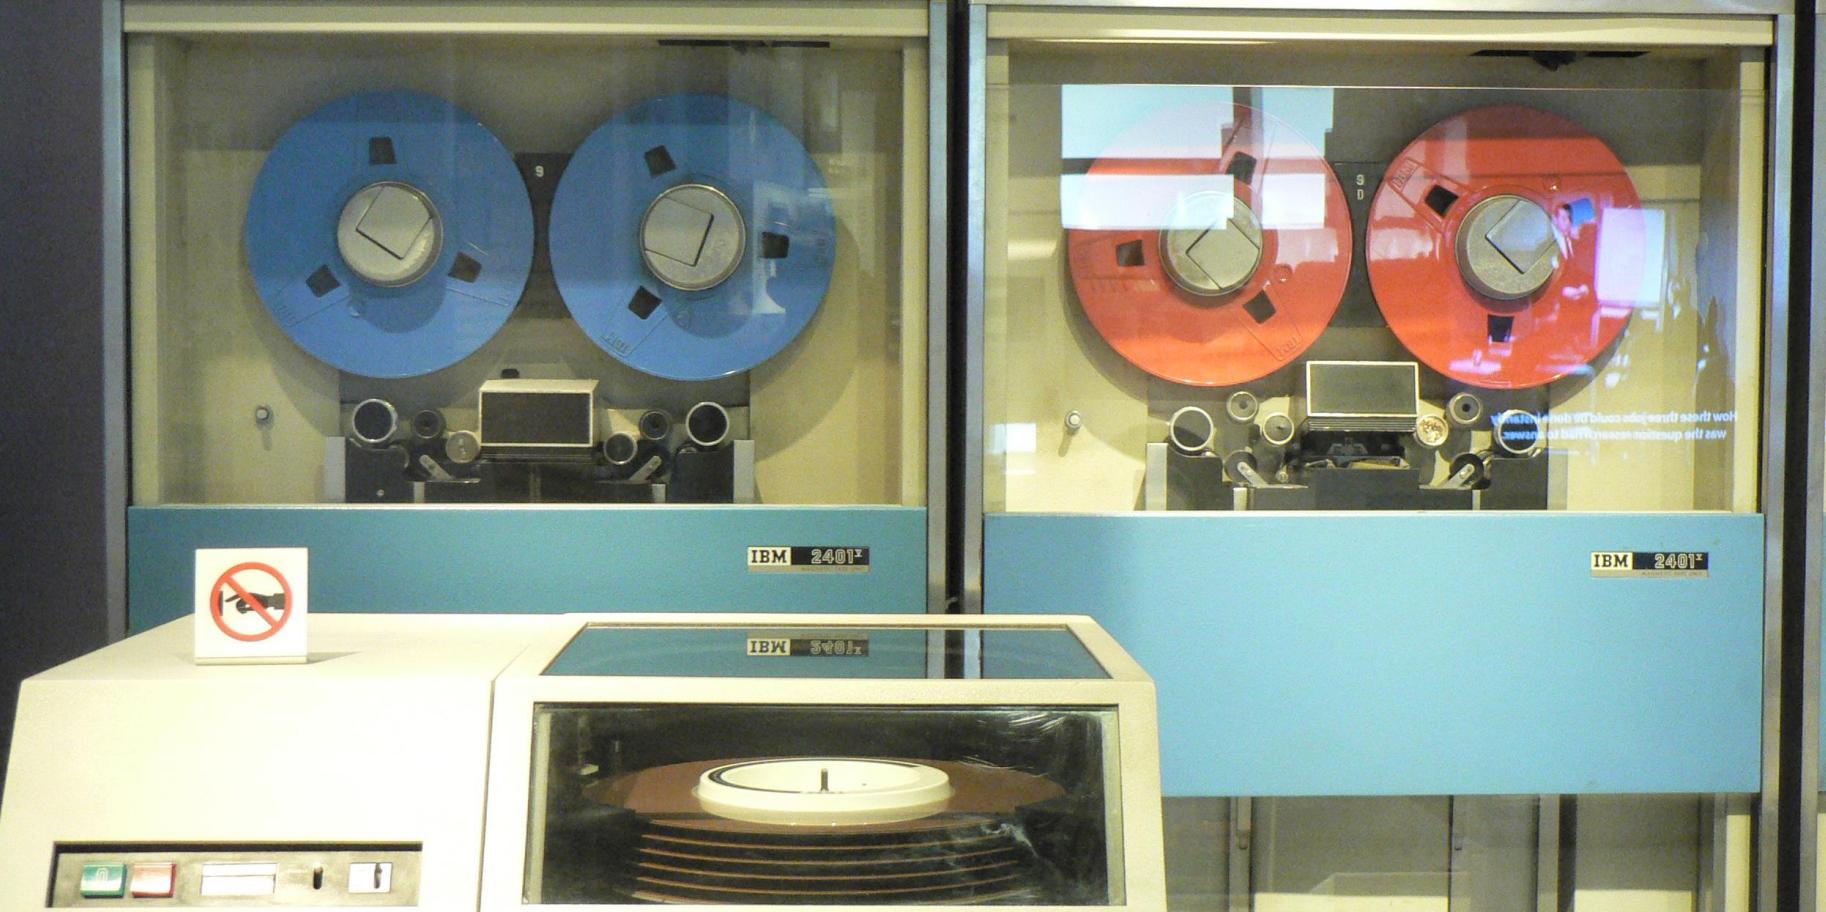
*shoppingcart-gamma.example.com*

<service-name>.<environment>.domain.com

*shoppingcart.gamma.example.com*

# Use a Dynamic Service Registry

- Avoids the DNS TTL issue
- More than service registry & discovery
  - Configuration management
  - Health checks
- Plenty of options
  - ZooKeeper (Apache)
  - Eureka (Netflix)
  - Consul (HashiCorp)
  - SmartStack (Airbnb)

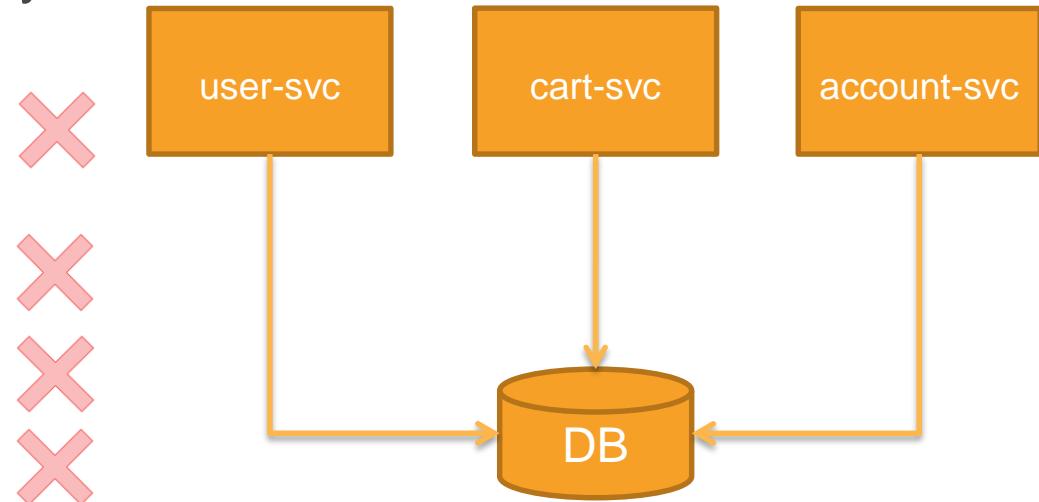


Data management

# Challenge: Centralized Database

Monolithic applications typically have a monolithic data store:

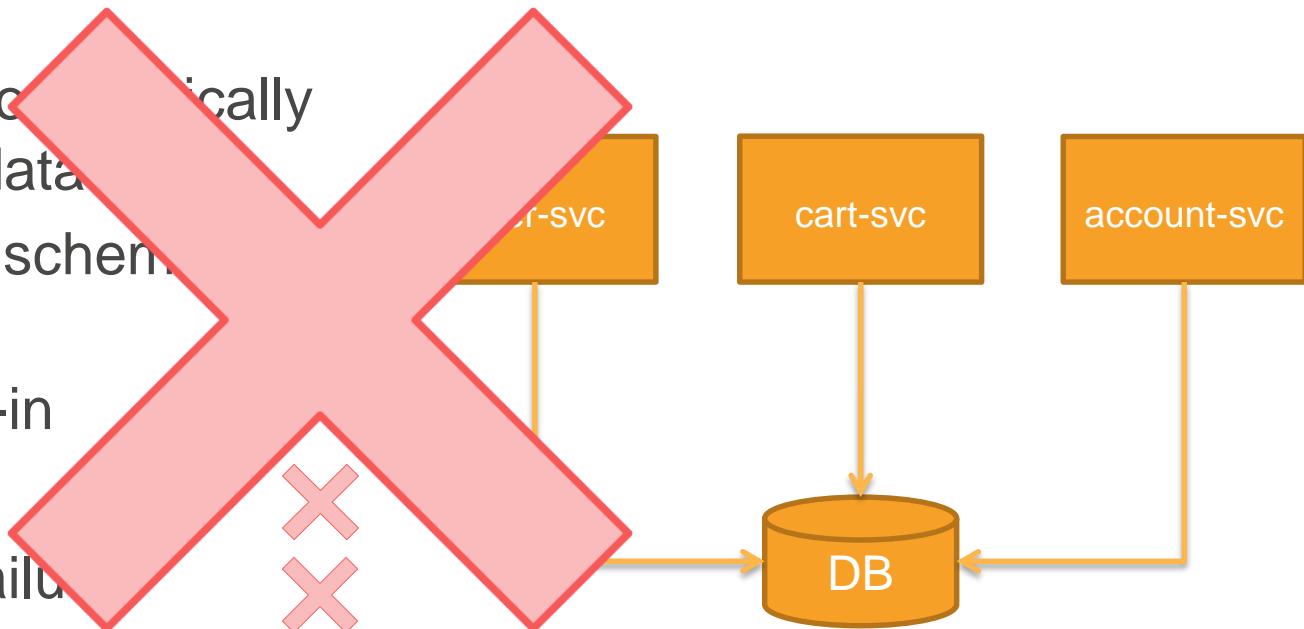
- Difficult to make schema changes
- Technology lock-in
- Vertical scaling
- Single point of failure



# Centralized Database – Anti-pattern

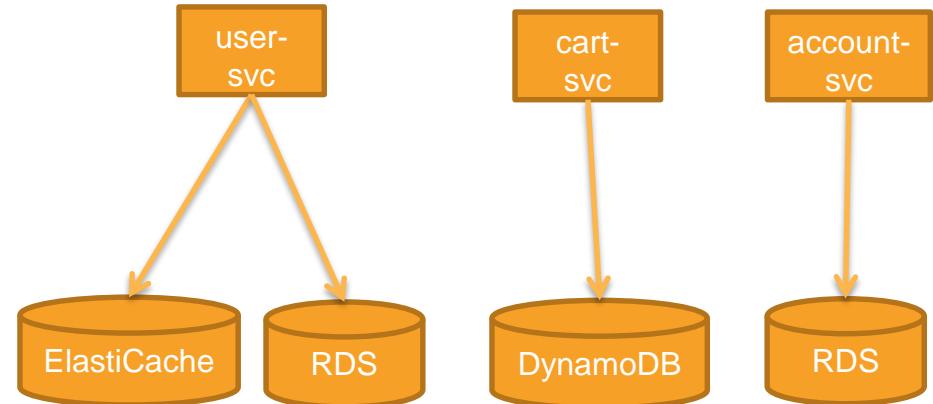
Monolithic applications typically have a monolithic database.

- Difficult to make schema changes
- Technology lock-in
- Vertical scaling
- Single point of failure



# Decentralized Data Stores

- Each service chooses its data store technologies
- Low impact schema changes
- Independent scalability
- Data is gated through the service API



# Challenge: Transactional Integrity

- Use a pessimistic model
  - Handle it in the client
  - Add a transaction manager / distributed locking service
  - Rethink your design
- Use an optimistic model
  - Accept eventual consistency
  - Retry (if idempotent)
  - Fix it later
  - Write it off

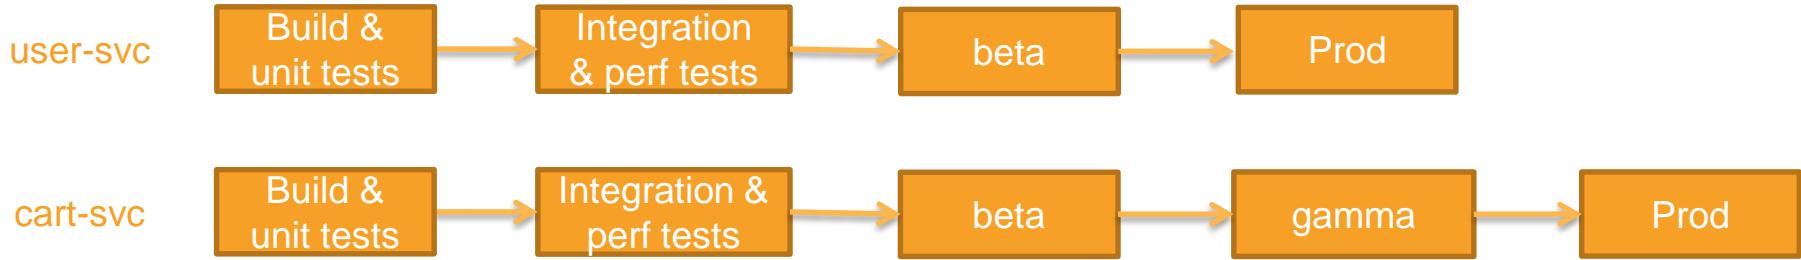
# Challenge: Aggregation

- **Pull:** Make the data available via your service API
- **Push:** To Amazon S3, Amazon CloudWatch, or another service you create
- **Pub/sub:** Via Amazon Kinesis or Amazon SQS



Deployment

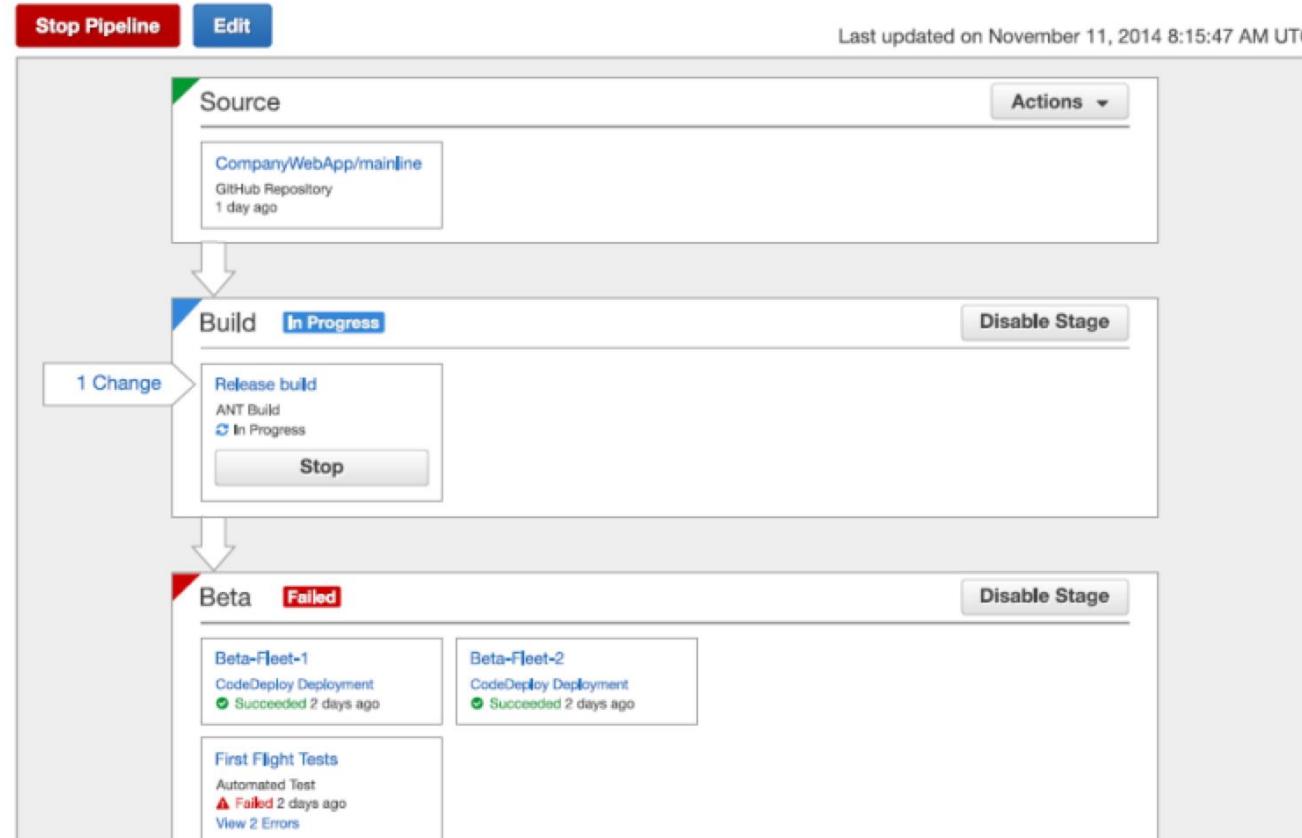
# Continuous Delivery & Continuous Deployment



Create the right build pipeline for each service

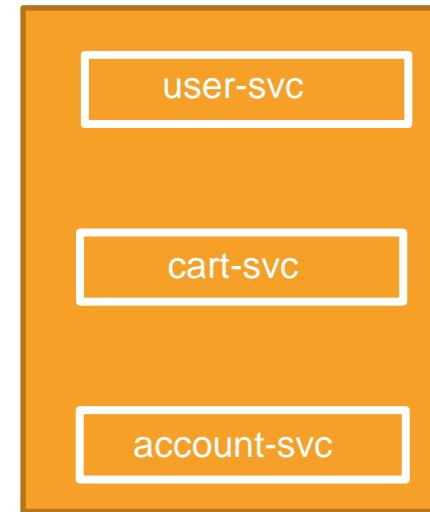
- AWS CodeDeploy
- AWS Elastic Beanstalk
- Jenkins, CircleCI, Travis,...

# AWS CodePipeline



# Multiple Services per Container/Instance

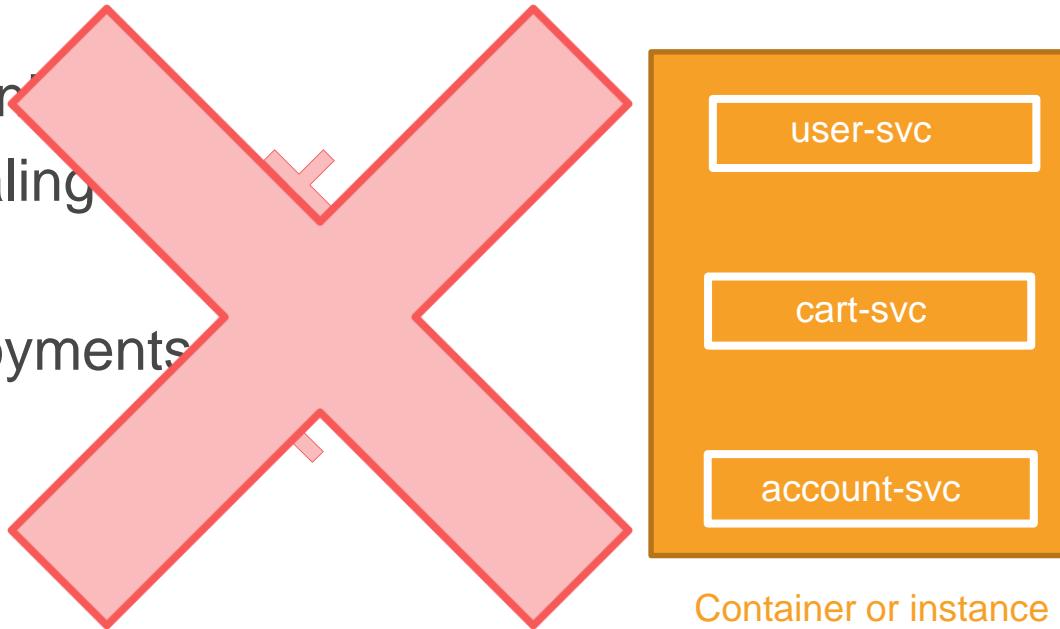
- Independent monitoring
- Independent scaling
- Clear ownership
- Immutable deployments



Container or instance

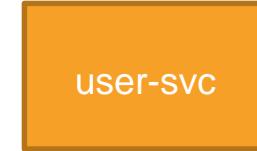
# Multiple Services per Container/Instance – Anti-pattern

- Independent monitoring
- Independent scaling
- Clear ownership
- Immutable deployments

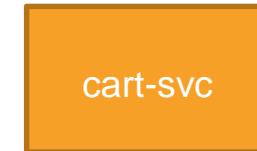


# Single Service per Container/Instance

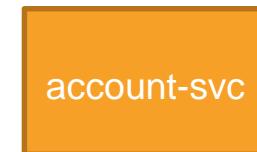
- Independent monitoring
- Independent scaling
- Clear ownership
- Immutable deployments



container or instance

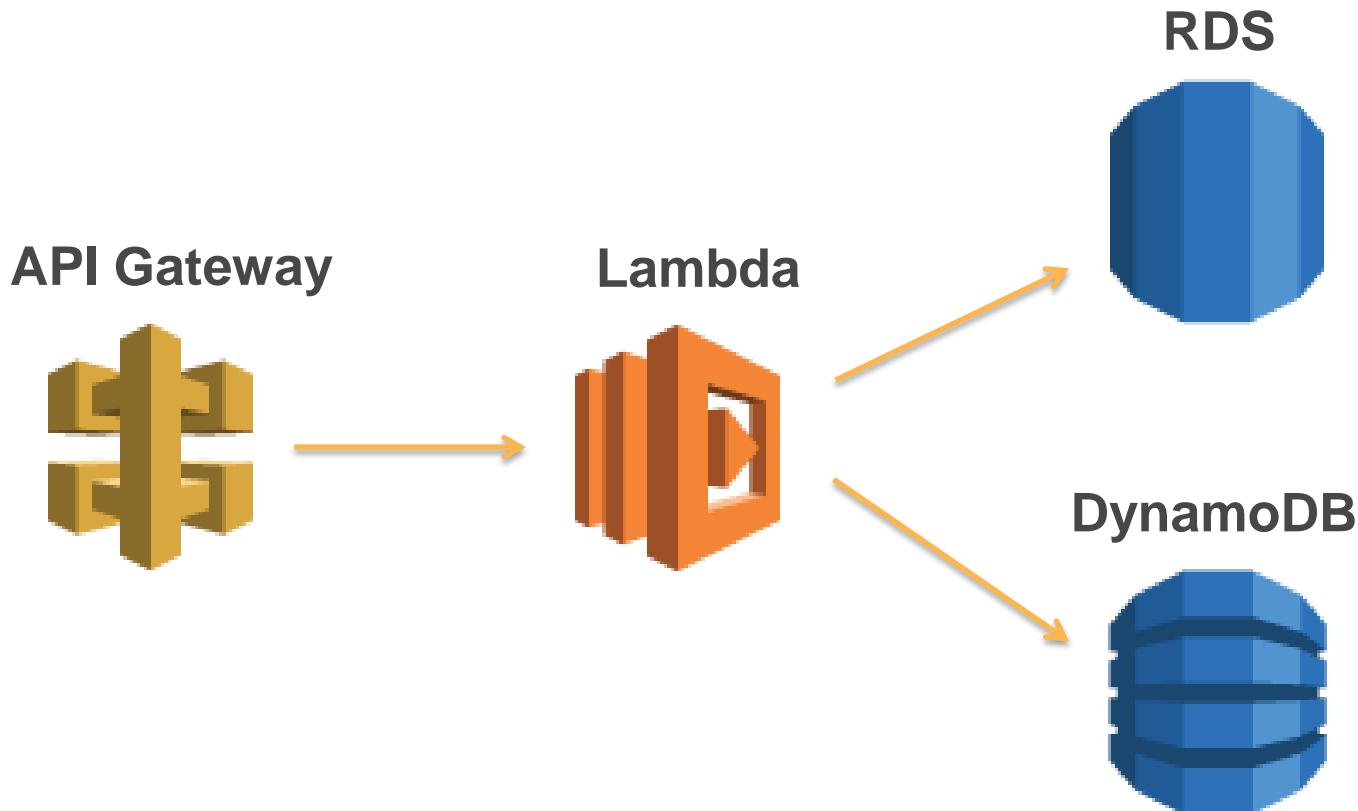


container or instance



container or instance

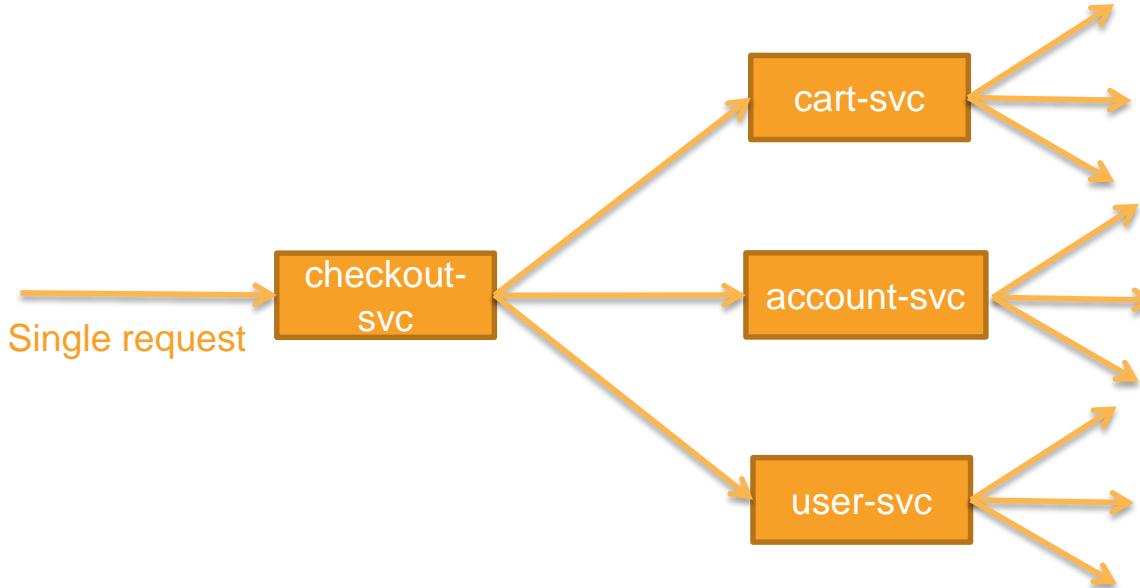
# ...Or Just Use AWS Lambda



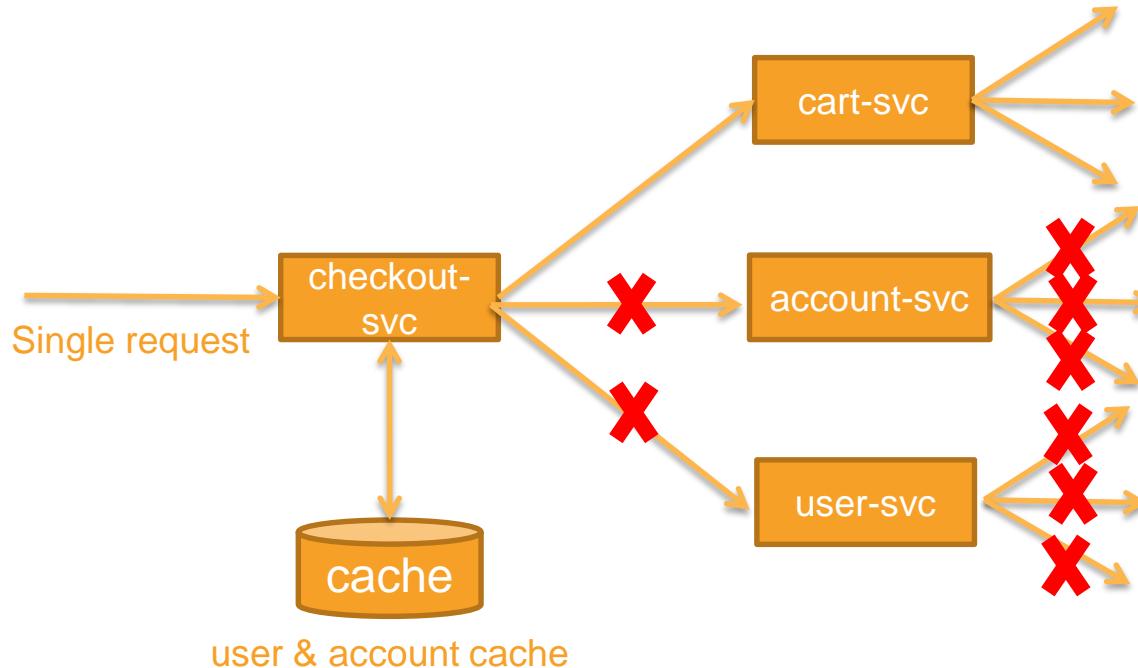


I/O explosion

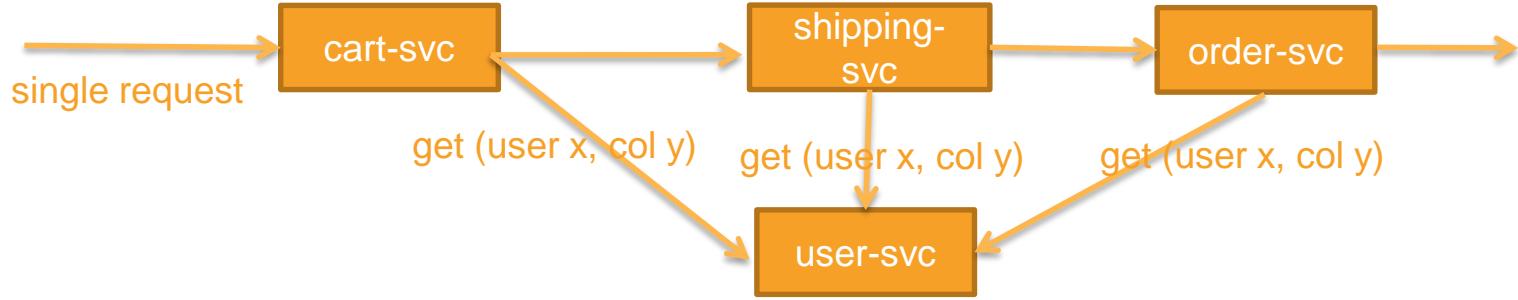
# Challenge: Request Multiplication



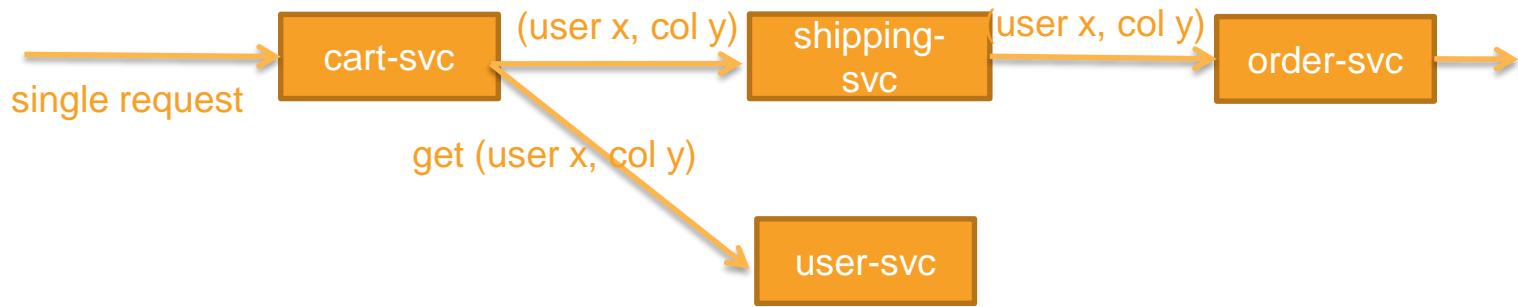
# Add Client Caching

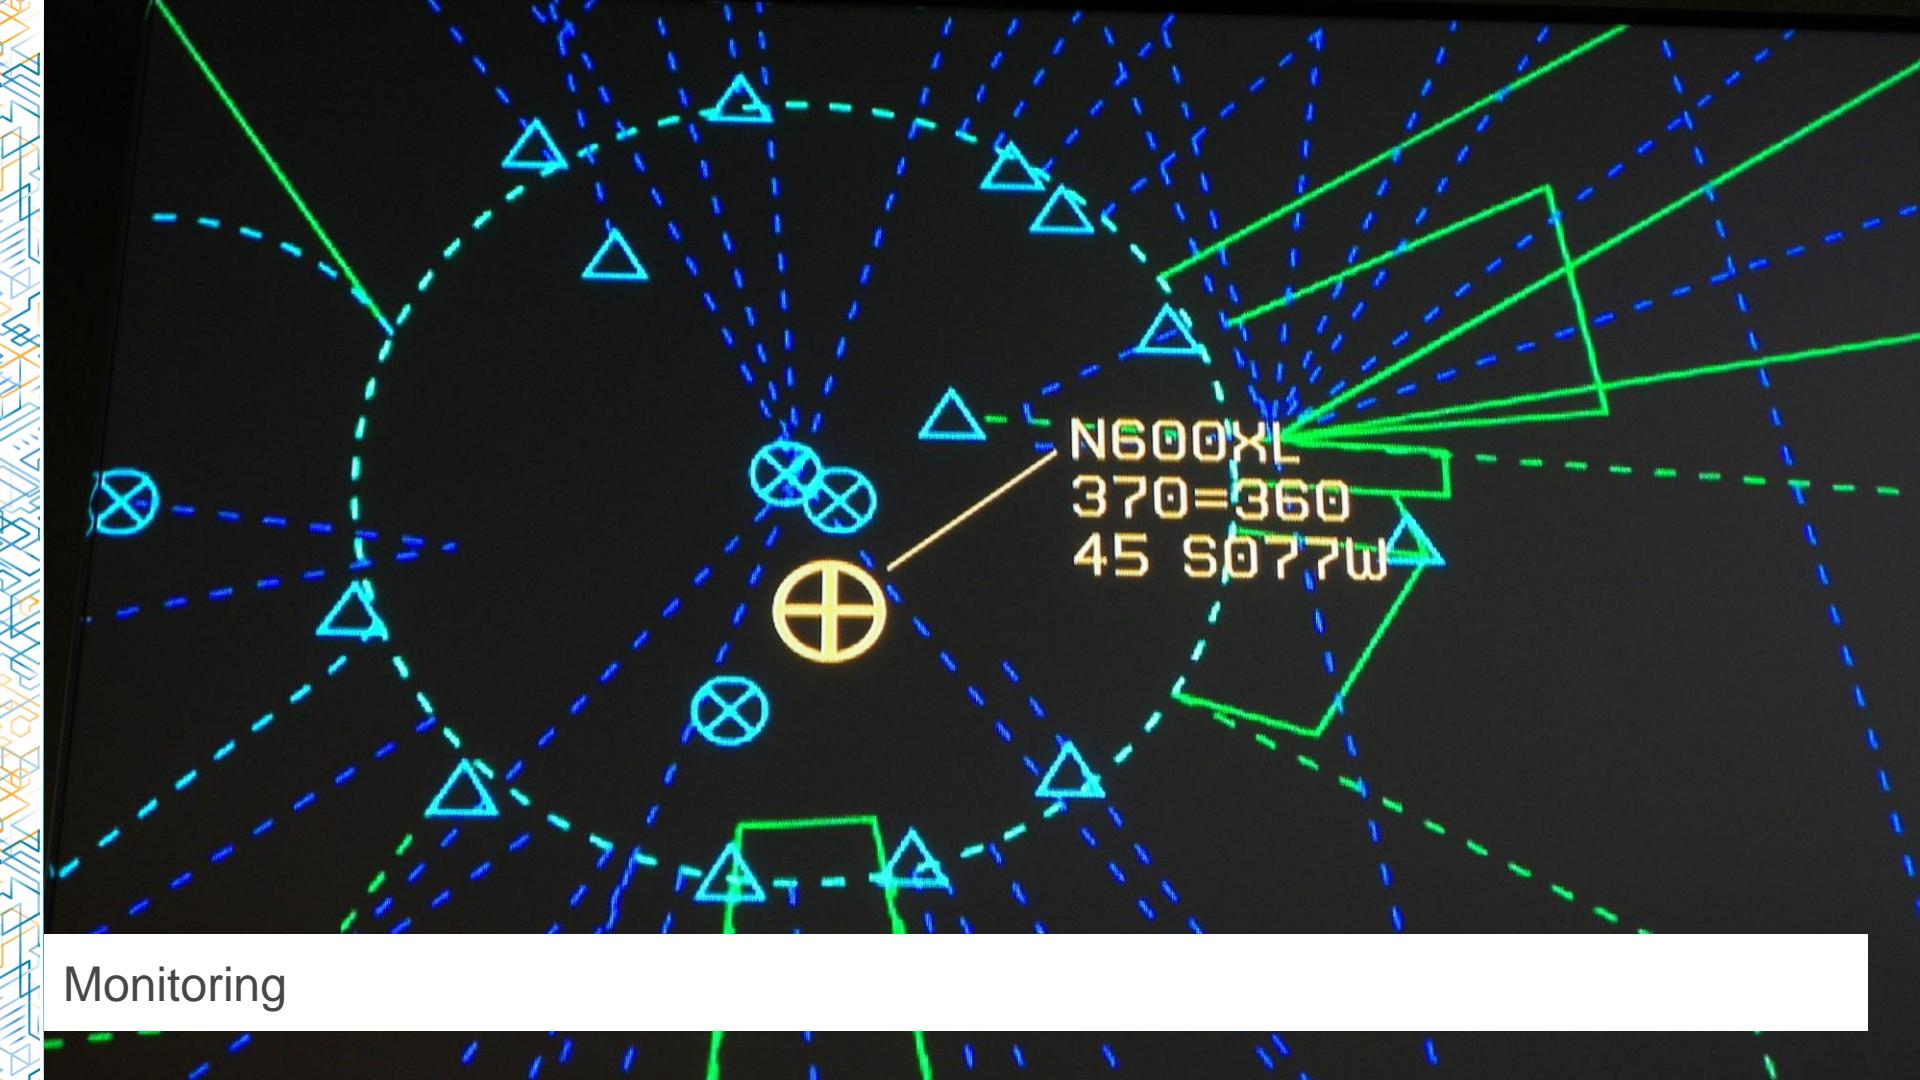


# Challenge: Hotspots



# Use Dependency Injection





N600XL  
370=360  
45 S077W

Monitoring

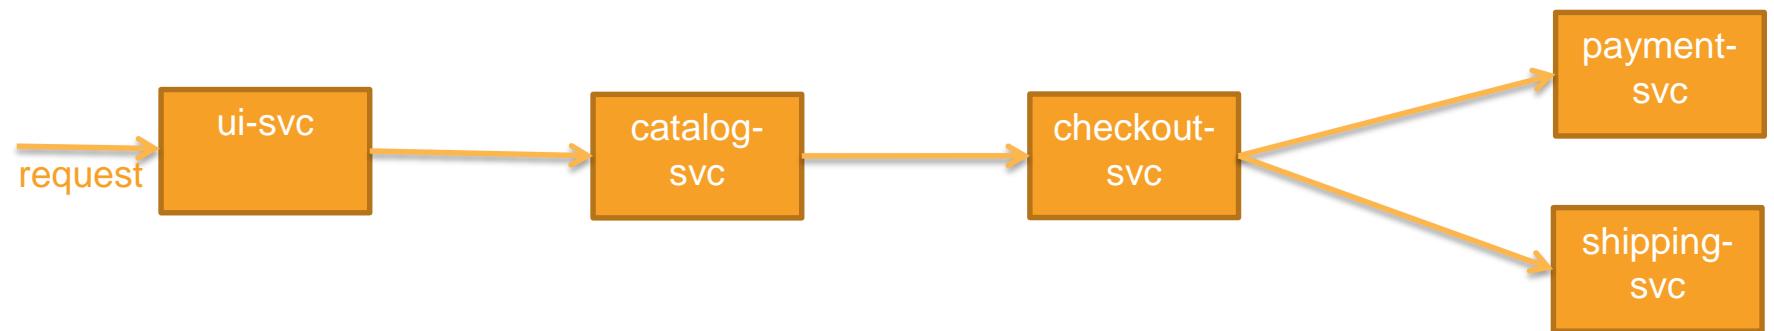
# Challenge: monitoring

- *Publish* externally relevant metrics
  - Latency
  - RPS
  - Error rate
- *Understand* internally relevant metrics
  - Basic – CloudWatch
  - OS
  - Application

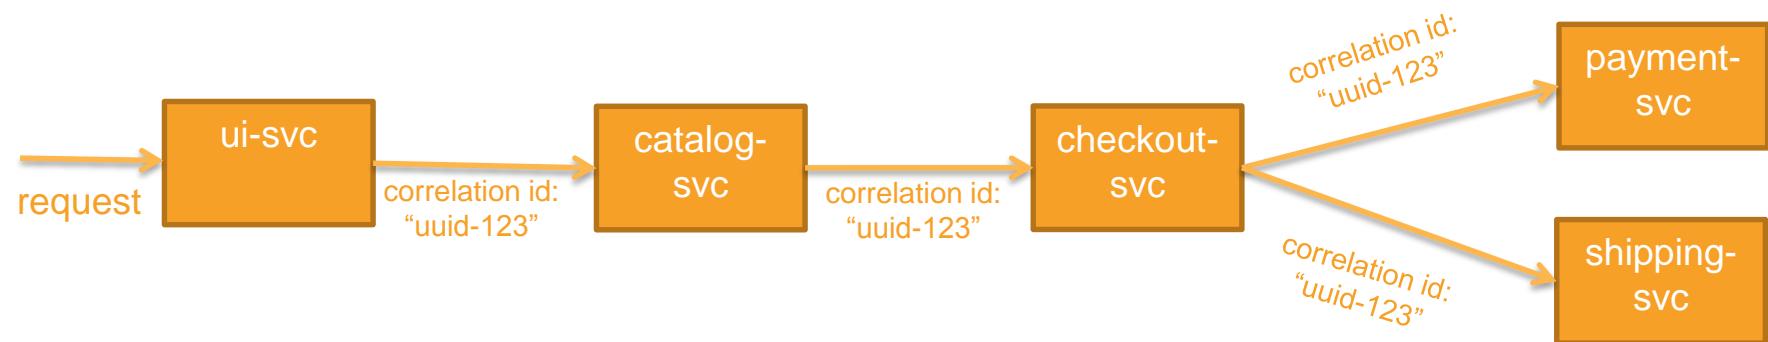
# Challenge: Logging

- Pick a common log aggregation solution
- Agree on log entry formats
- Use naming conventions
- Agree on correlation strategy

# Challenge: Correlating Requests



# Use Correlation IDs



```
09-02-2015 15:03:24 ui-svc INFO [uuid-123] ....  
09-02-2015 15:03:25 catalog-svc INFO [uuid-123] ....  
09-02-2015 15:03:26 checkout-svc ERROR [uuid-123] ....  
09-02-2015 15:03:27 payment-svc INFO [uuid-123] ....  
09-02-2015 15:03:27 shipping-svc INFO [uuid-123] ....
```

# What did we cover?

- Ownership
- Discovery
- Data management
- Deployment
- I/O explosion
- Monitoring



# Related Sessions

- ARC201 - Microservices Architecture for Digital Platforms with AWS Lambda, Amazon CloudFront and Amazon DynamoDB
- CMP302 - Amazon EC2 Container Service: Distributed Applications at Scale
- DEV203 - Using Amazon API Gateway with AWS Lambda to Build Secure and Scalable APIs
- DVO401 - Deep Dive into Blue/Green Deployments on AWS
- SPOT304 - Faster, Cheaper, Safer Products with AWS: Adrian Cockcroft Shares Experiences Helping Customers Move to the Cloud



**Remember to complete  
your evaluations!**



# Thank you!

Adrian Trenaman

SVP Engineering, gilt.com

@adrian\_trenaman

Derek Chiles

Sr. Mgr, Solutions Architecture, AWS

@derekchiles